



Introduction to LLVM

Sthiti Deka

Jayashree Venkatesh

Agenda

Time	Elapsed Time	Topic	Description
5 min.	5 min.	Welcome, Course Agenda and Introductions	Welcome participants Cover Goals and Objectives Review Agenda, Logistics & Take-away
10 min.	15 min.	What is LLVM & Why LLVM	Define the scope and motivation behind LLVM
45 min.	60 min.	LLVM Compiler Pipeline Overview	Provide a design overview of the LLVM Compiler Pipeline stages – the frontend, the mid-end and the back-end
15 min.	75 min.	Hands-On Ex1: Hello World in Clang	Learn how to compile, diagnose errors with clang, lli, llc and the command line options
10 min.	85 min.	Hands-On Ex2: LLVM C API to generate an LLVM IR Module	Learn how to use the LLVM API and to generate a module in LLVM IR
5 min.	90 min.	Hands-On Ex3: Print & View Program State & Flow	Learn how to use LLVM tools to view & print program flow and program states
10 min.	100 min.	Hands-On Ex4: Analysis Pass in LLVM API	Learn how to write an analysis pass in LLVM API
5min.	105 min.	Hands-On Ex5: Transformation Pass in LLVM API	Learn how to write a transformation pass in LLVM API
5 min.	110 min.	Hands-On Ex6: Debug tools in LLVM IR	Learn how to use print before and after transformation to debug
10 min	120 min.	Q/A and wrap up	Final questions and conclusions

Content

- What is LLVM?
- Why LLVM?
- LLVM Compiler Pipeline
 - LLVM Front-end
 - LLVM Mid-end
 - LLVM Back-end
- Hands-On

What is LLVM?

What is LLVM

- Open Source Project

What is LLVM

- Open Source Project
 - Compiler Framework

What is LLVM

- Open Source Project
 - Compiler Framework
- Infrastructure

What is LLVM

- Open Source Project
 - Compiler Framework
 - Optimizations
- Infrastructure

What is LLVM

- Open Source Project
 - Compiler Framework
 - Optimizations
 - API + IR
- Infrastructure

What is LLVM

- Open Source Project
 - Compiler Framework
 - Optimizations
 - API + IR
 - Retargetable: m languages x n targets
 - Infrastructure

What is LLVM

- Open Source Project
 - Compiler Framework
 - Optimizations
 - API + IR
 - Retargetable: m languages x n targets
 - Infrastructure
 - Reusable libraries
 - Tool Chain

Why LLVM?

Why LLVM

Why LLVM

- Reduced development time

Why LLVM

- Reduced development time
- Performance
 - Speed
 - Low Memory Footprint

Why LLVM


- Reduced development time
- Performance
 - Speed
 - Low Memory Footprint
- Ease of use
 - Toolchain, Libraries

LLVM Compiler Framework

LLVM Compiler Pipeline



LLVM Compiler Pipeline




LLVM Compiler Pipeline




LLVM Front End

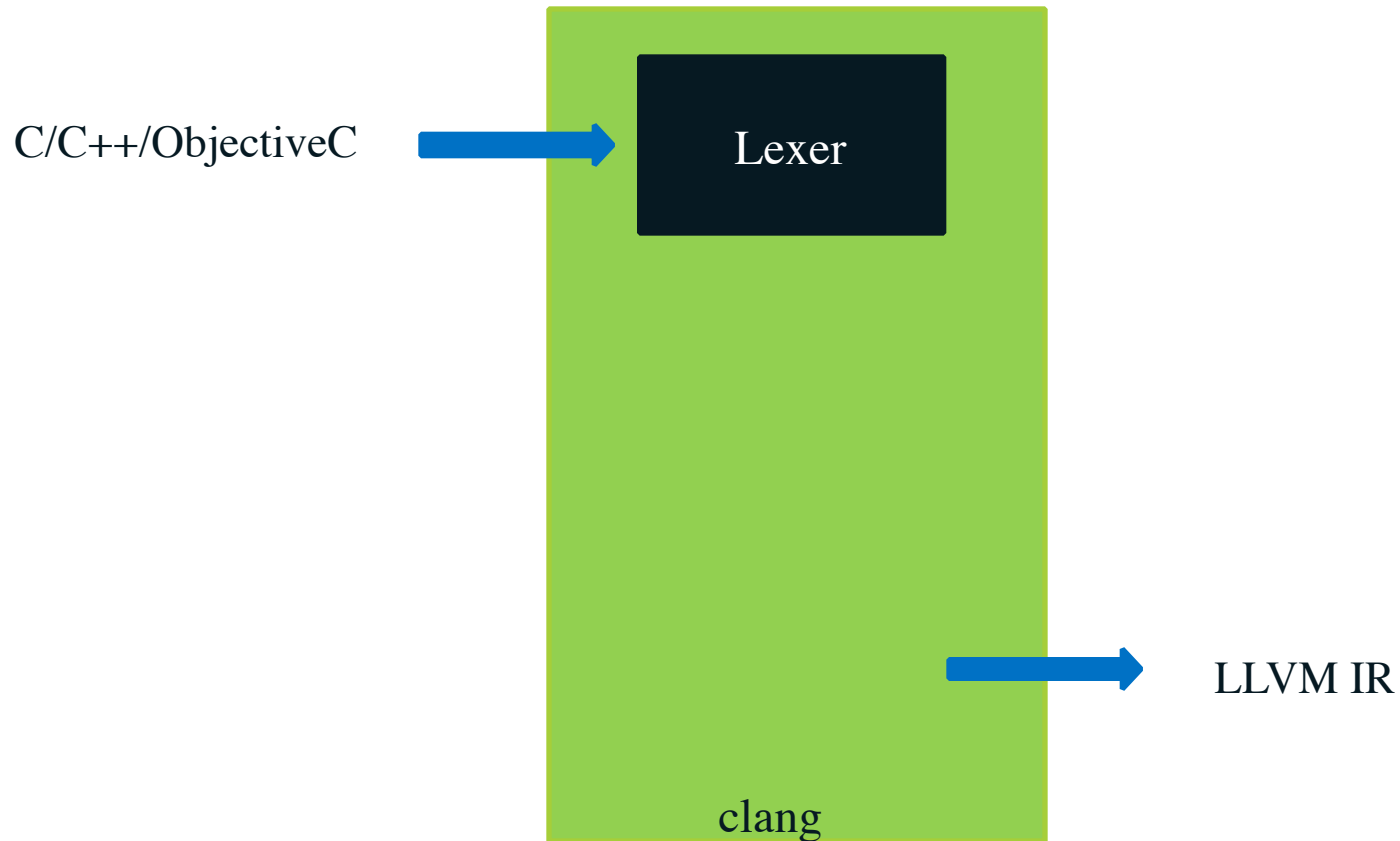
LLVM Pipeline



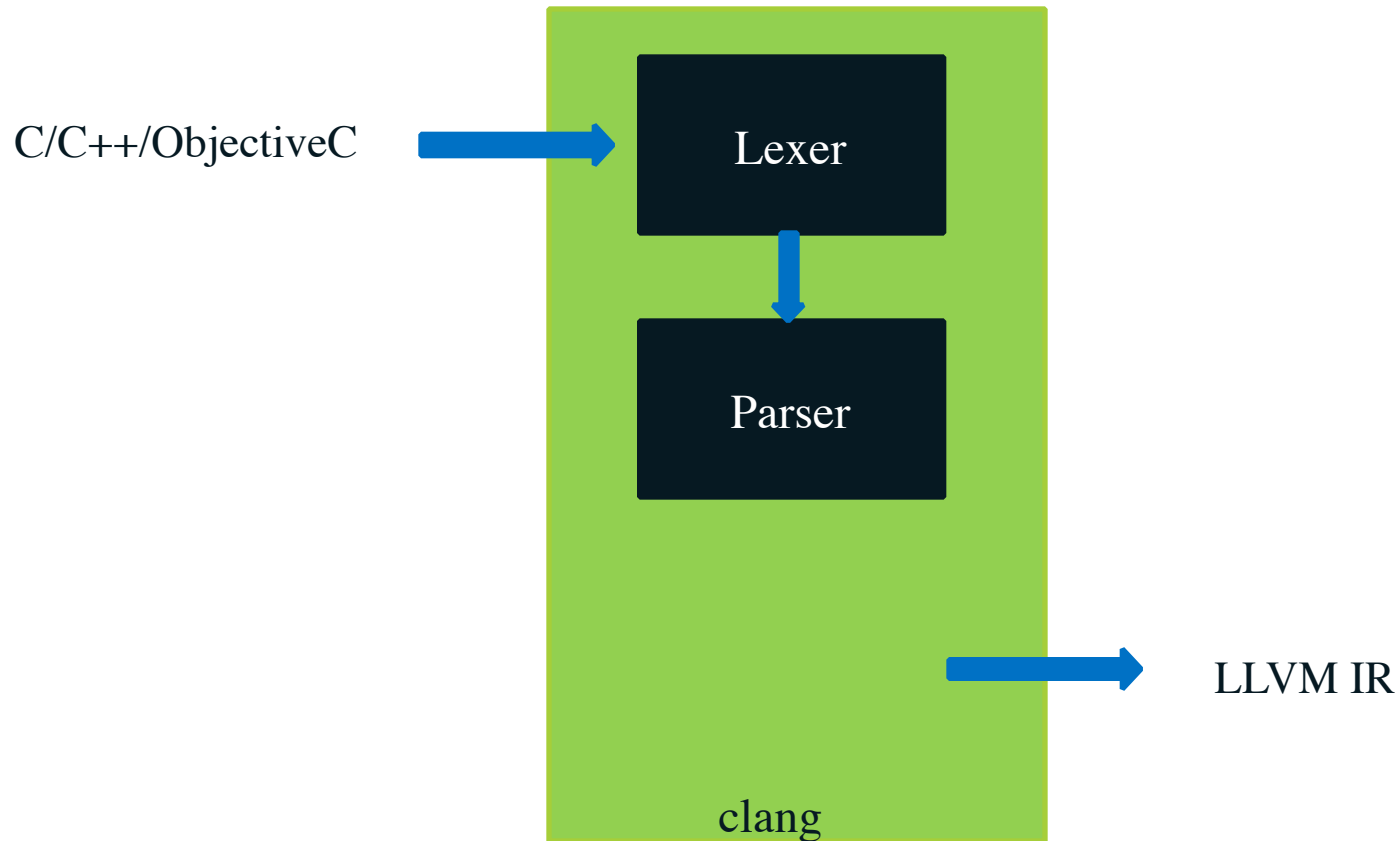
Clang – Design Overview



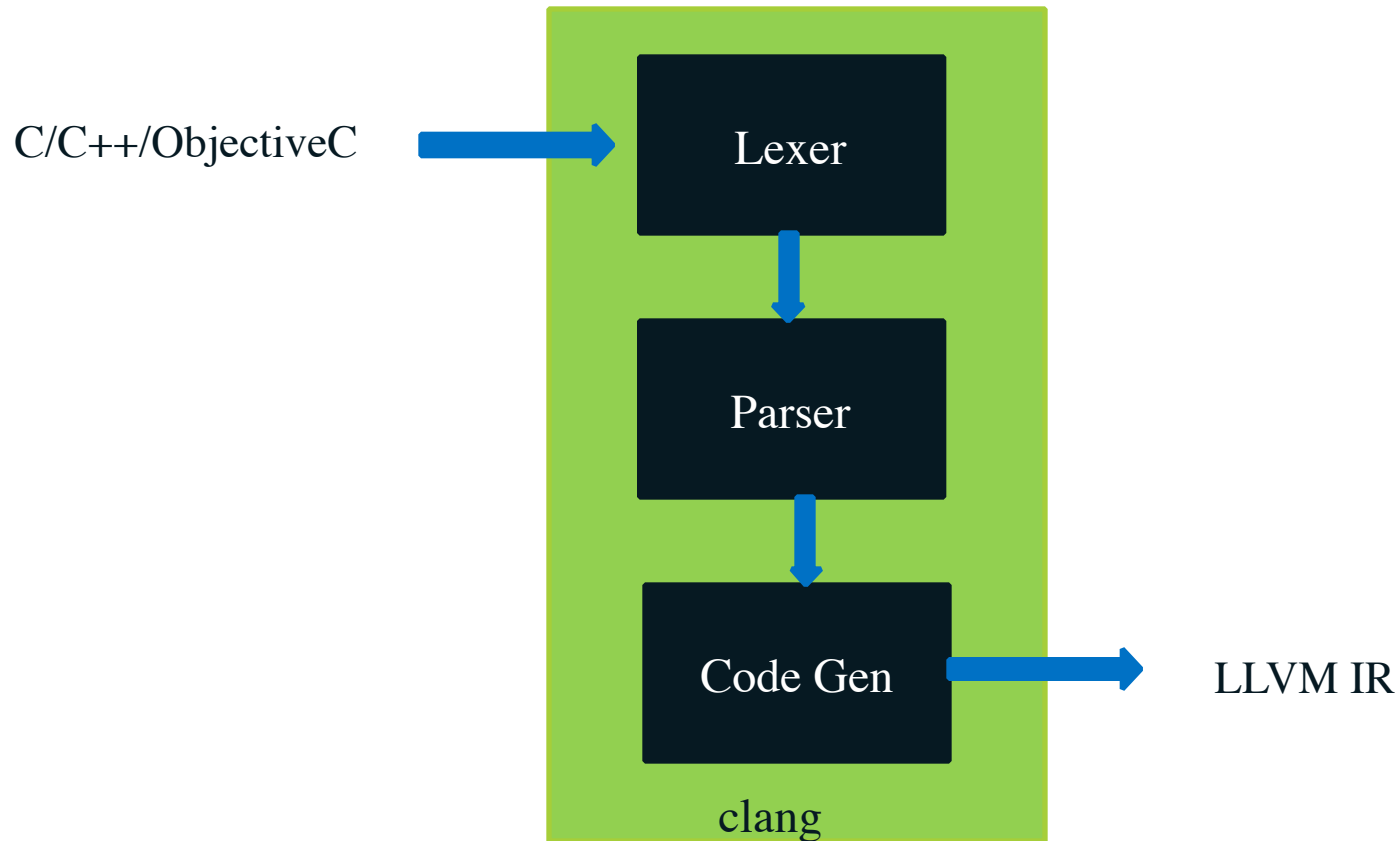
Clang – Design Overview



Clang – Design Overview



Clang – Design Overview



Example LLVM IR

```
int sum(int a, int b) {  
    return a+b;  
}
```

C function

Example LLVM IR

```
int sum(int a, int b) {  
    return a+b;  
}
```

C function

```
define i32 @sum(i32 %a, i32 %b)  
{  
entry:  
    %add = add i32 %b, %a  
    ret i32 %add  
}
```

LLVM IR function

Example LLVM IR SSA

```
int mul_add(  
int a,  
int b,  
int c)  
{  
    b = a * b;  
    b = c + b;  
    return b;  
}
```

C function

```
define i32 @mul_add(  
i32 %a,  
i32 %b,  
i32 %c)  
{  
entry:  
    %mul = mul i32 %b, %a  
    %add = add i32 %mul, %c  
    ret i32 %add  
}
```


LLVM IR function

LLVM Mid End Optimizer


LLVM Pipeline



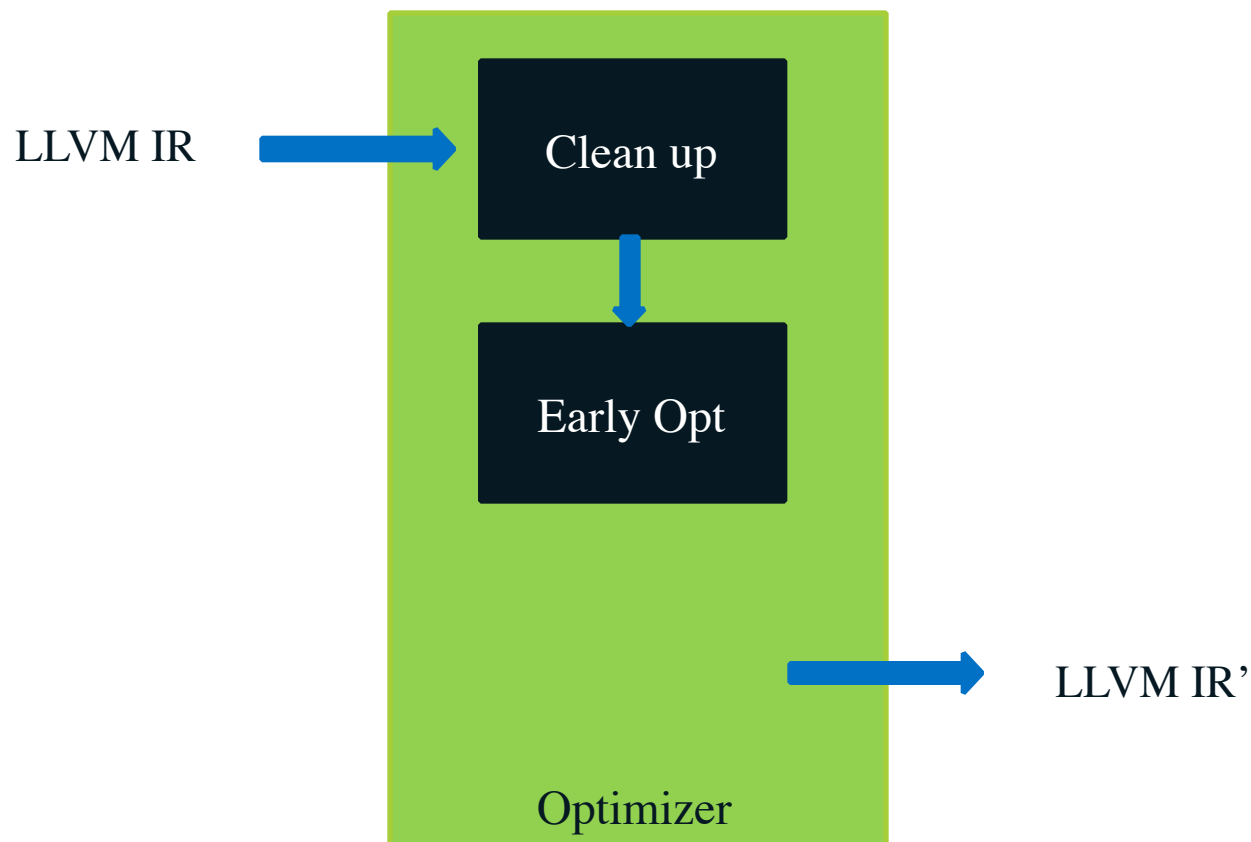
LLVM Mid End Optimizer – Design Overview



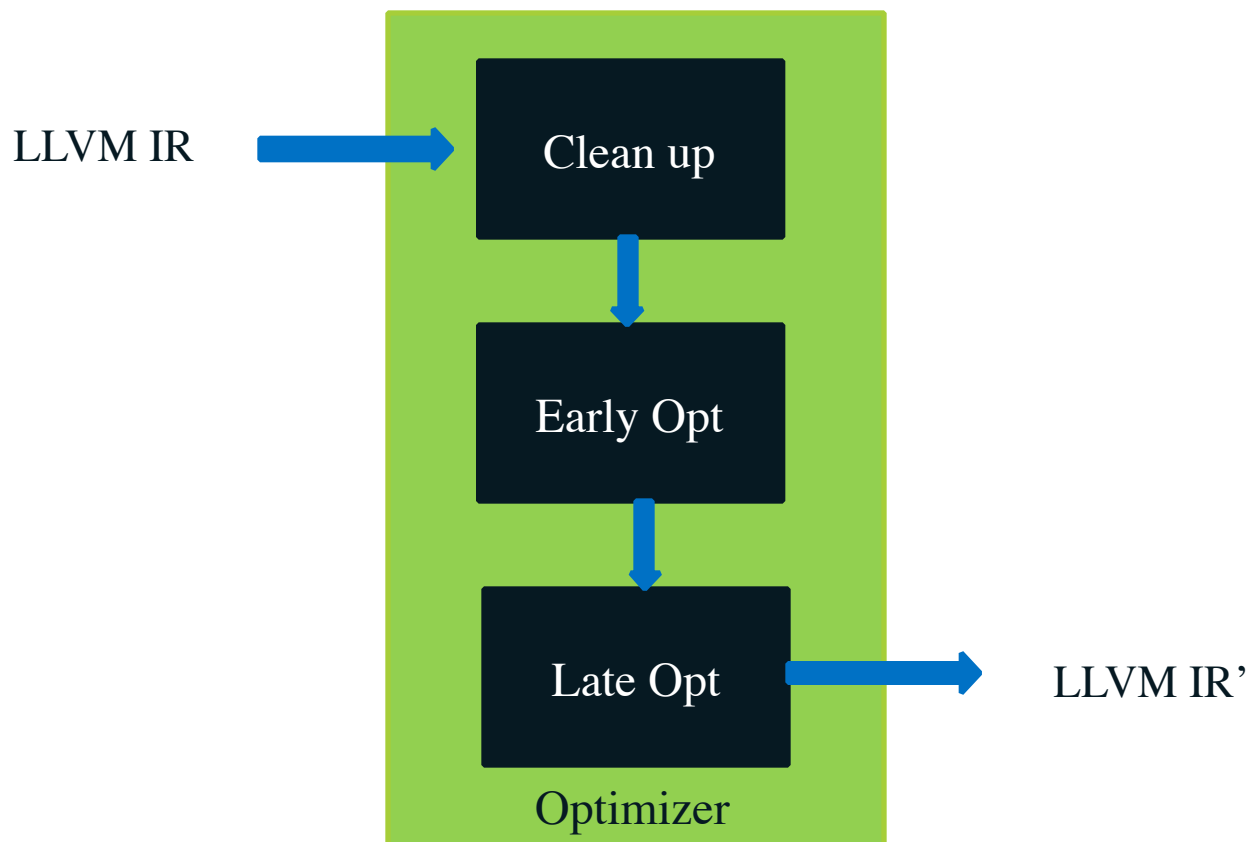
LLVM Mid End Optimizer – Design Overview



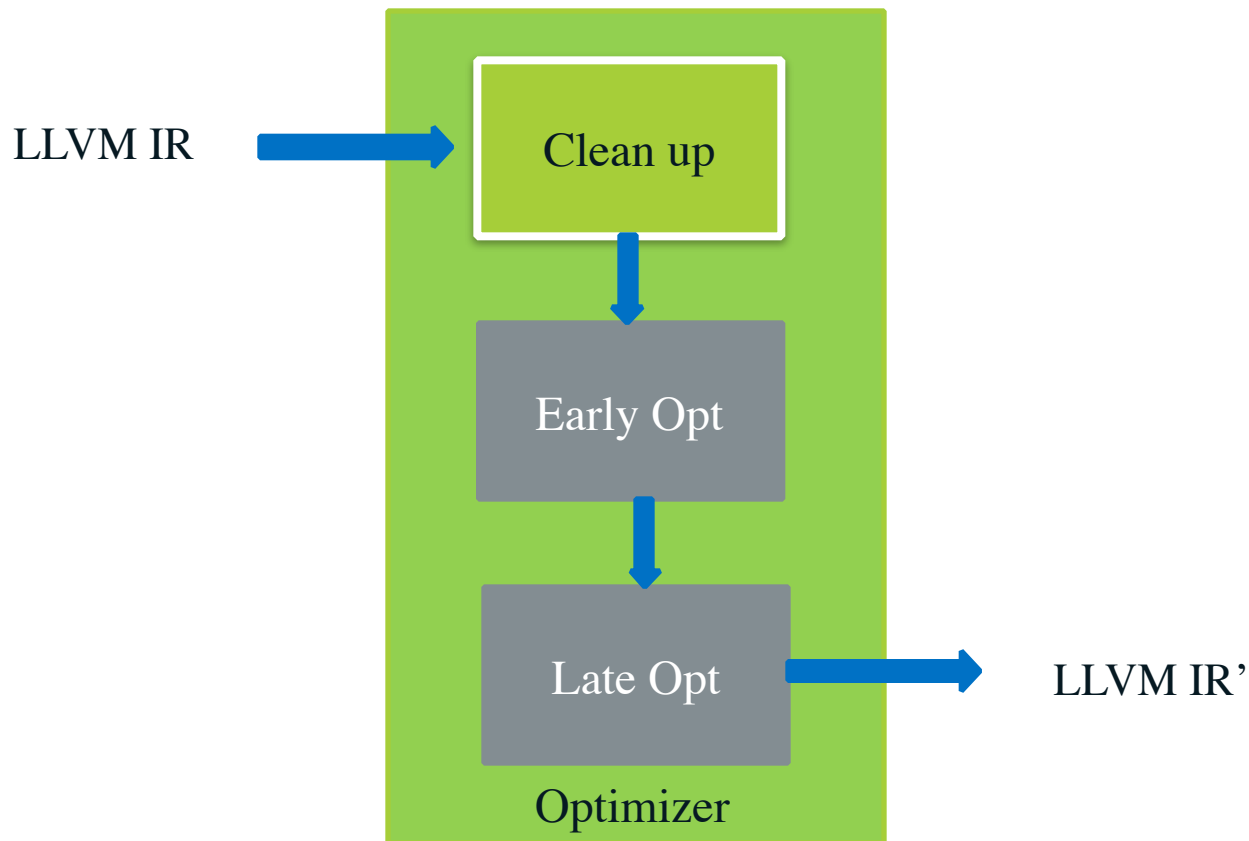
LLVM Mid End Optimizer – Design Overview



LLVM Mid End Optimizer – Design Overview



Mid End Optimizer – Clean up



Mid End Optimizer – Clean up

Front End
clang

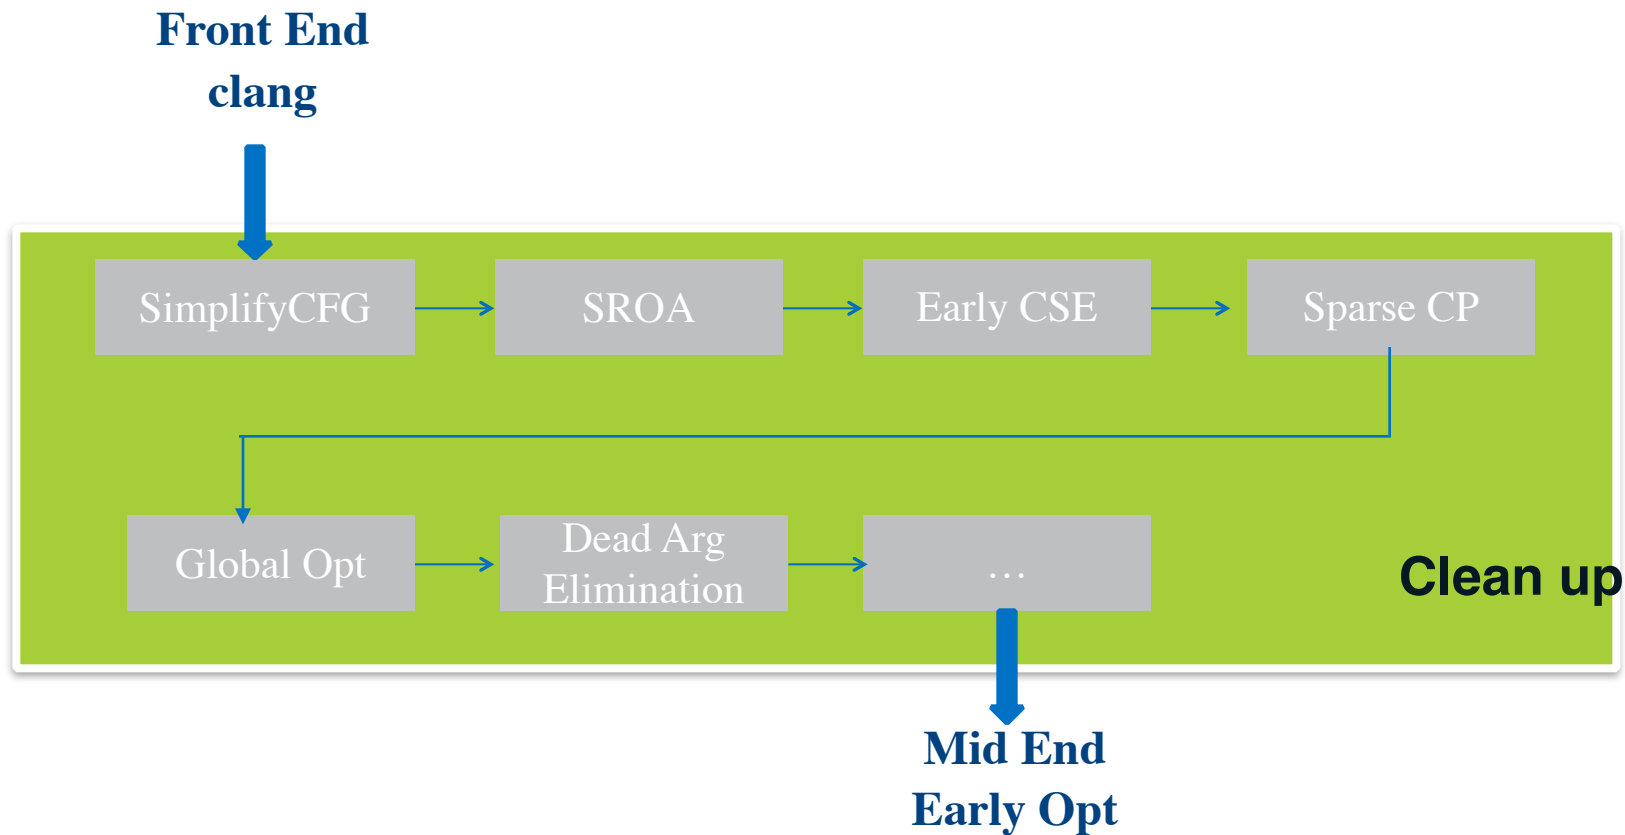


Clean up




Mid End
Early Opt

Mid End Optimizer – Clean up

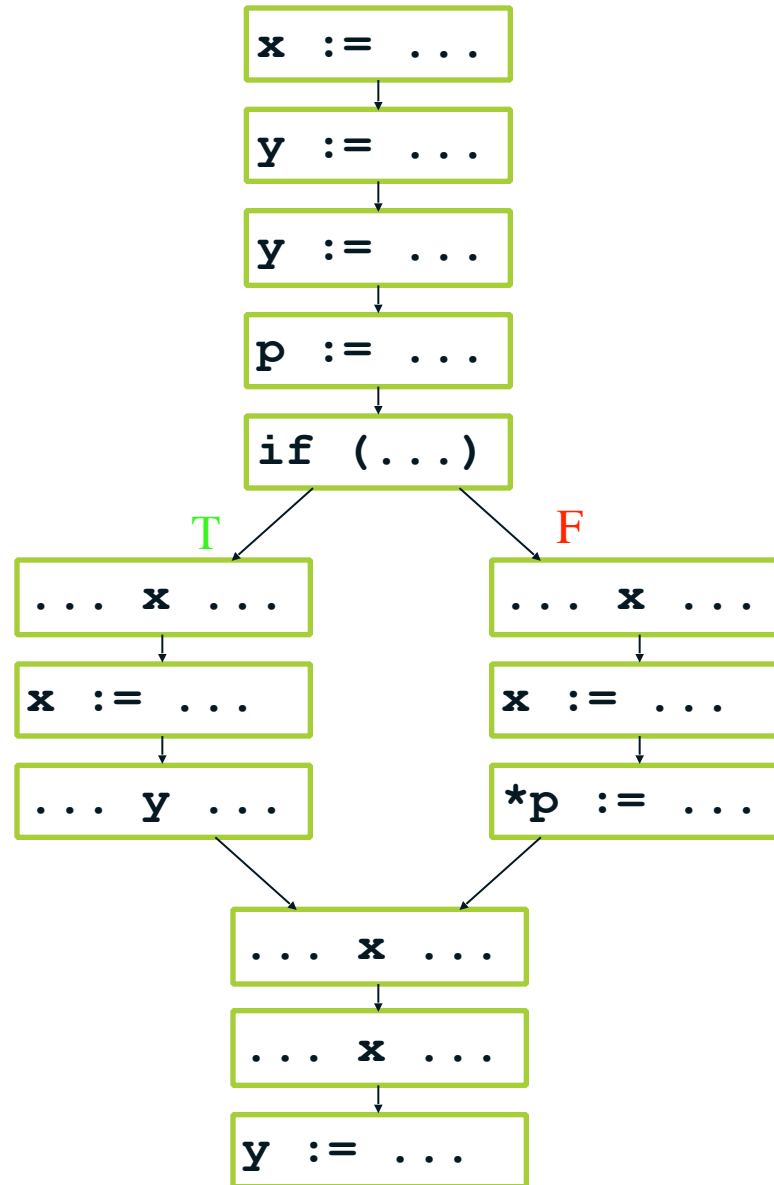


Mid End Optimizer – Clean up

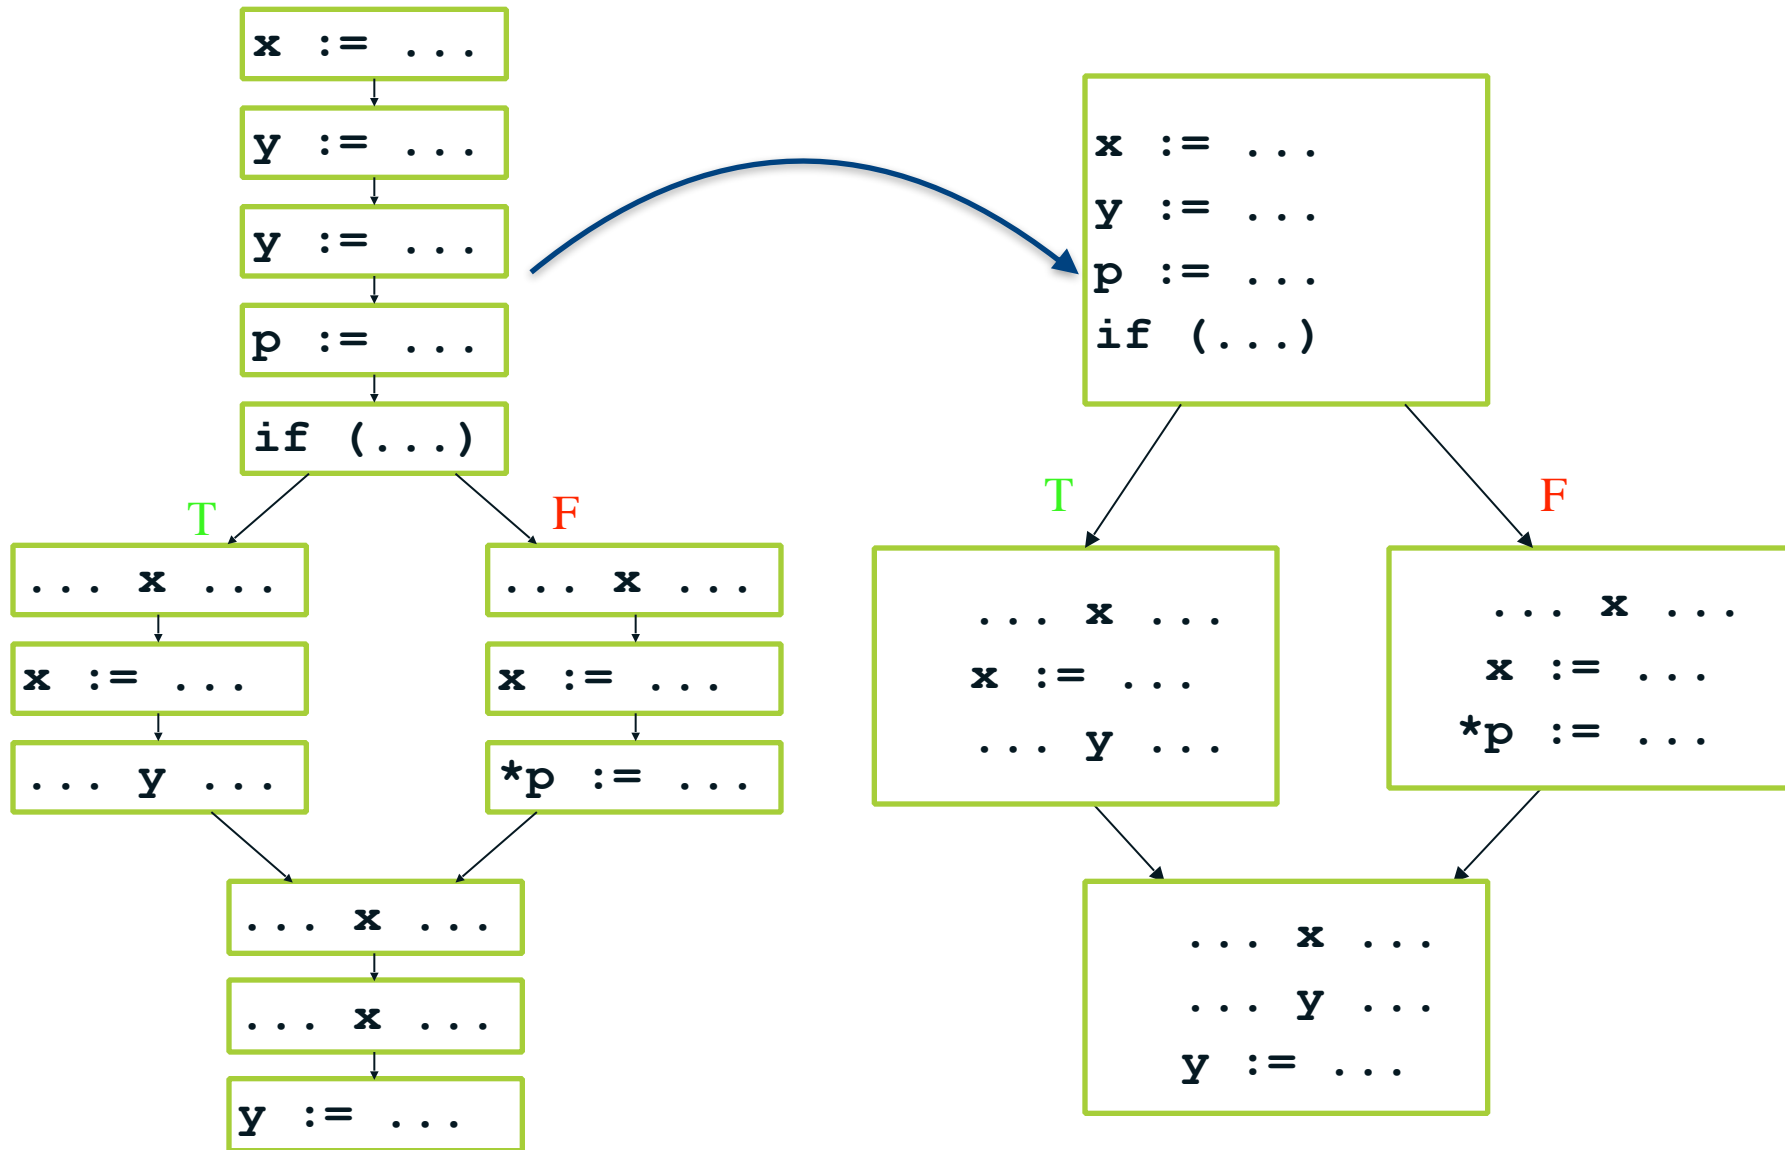


Example CFG

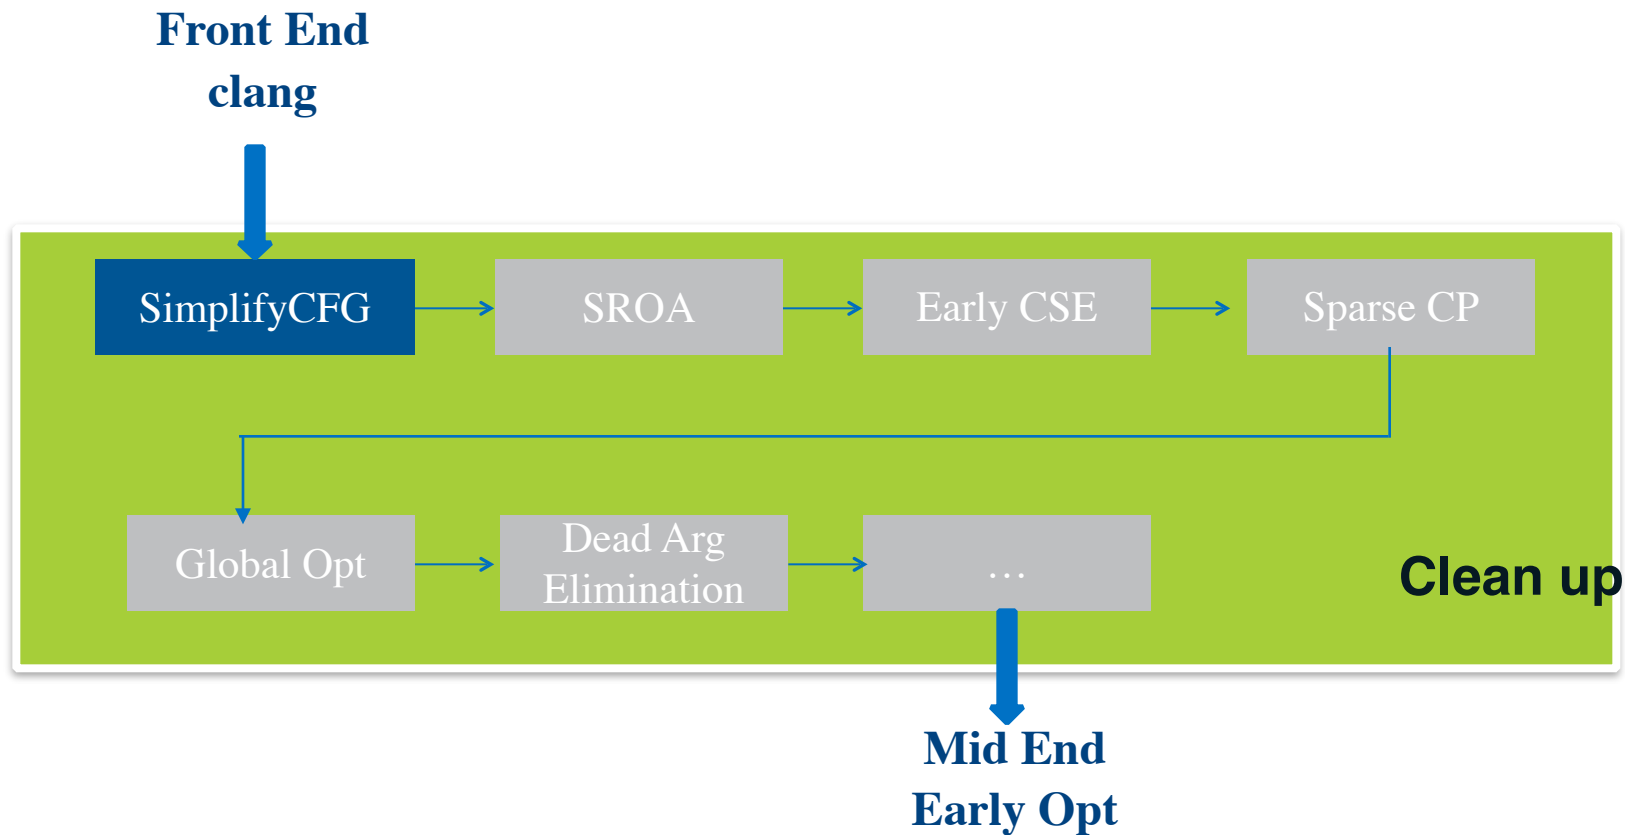
```
x := ...  
y := ...  
y := ...  
p := ...  
if (...) {  
    ... x ...  
    x := ...  
    ... y ...  
}  
else {  
    ... x ...  
    x := ...  
    *p := ...  
}  
... x ...  
... y ...  
y := ...
```



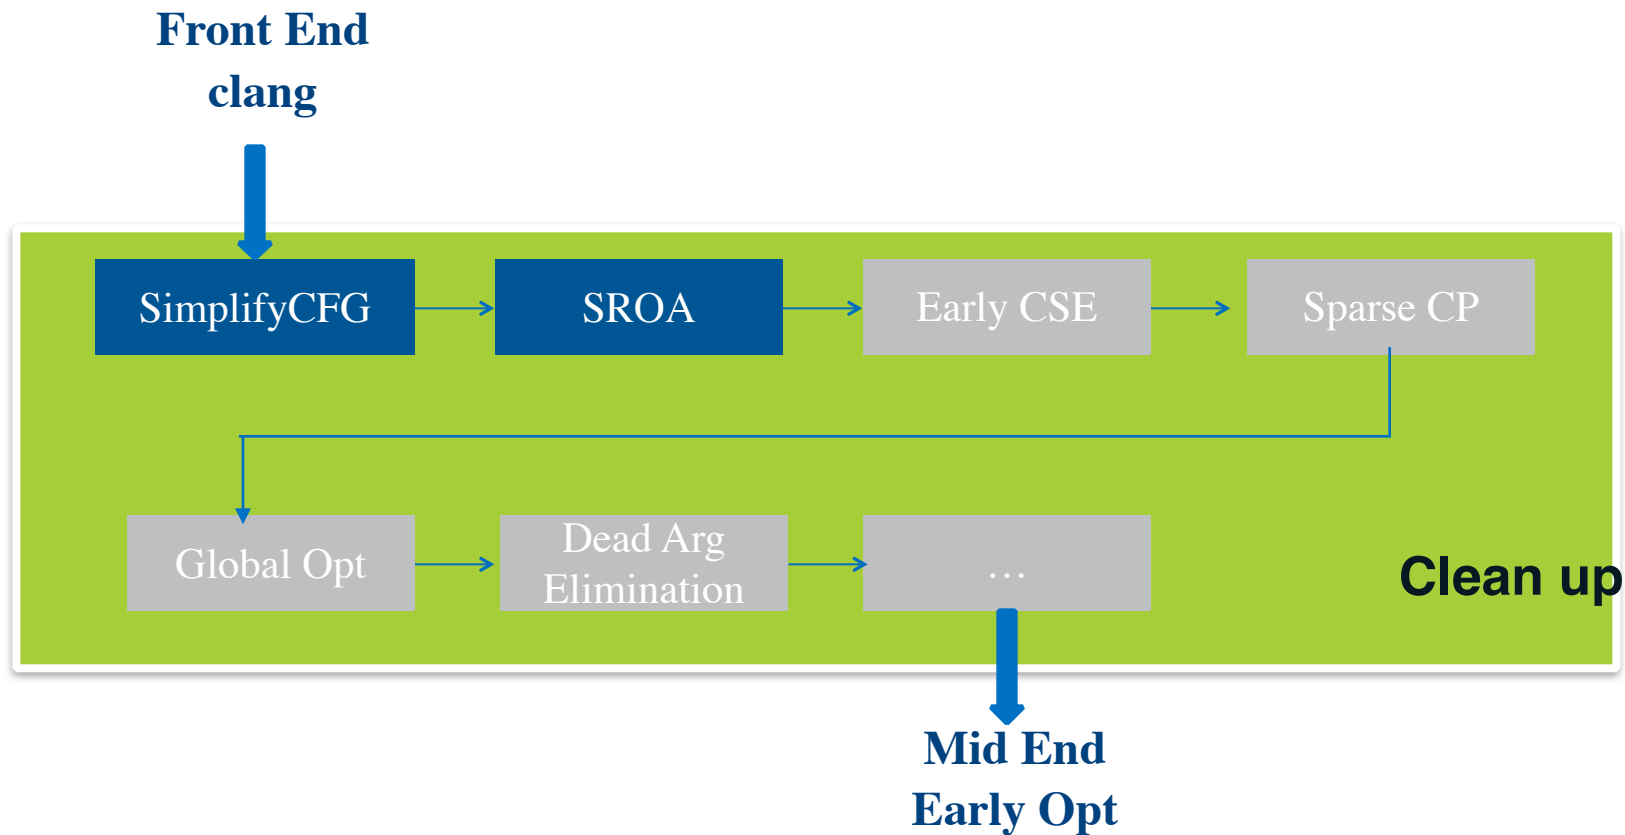
Simplified CFG



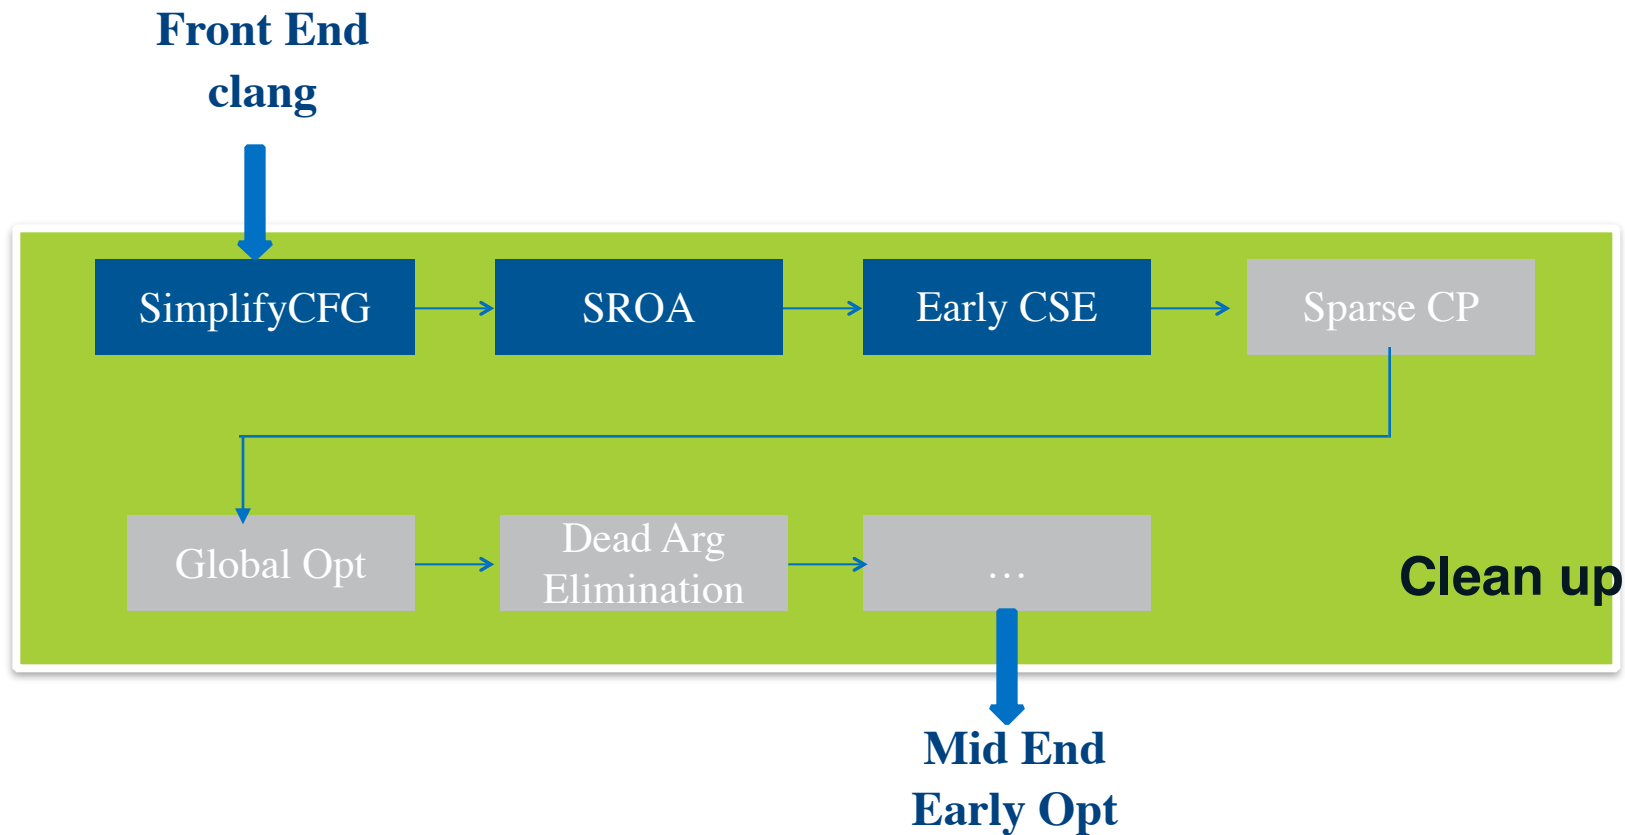
Mid End Optimizer – Clean up



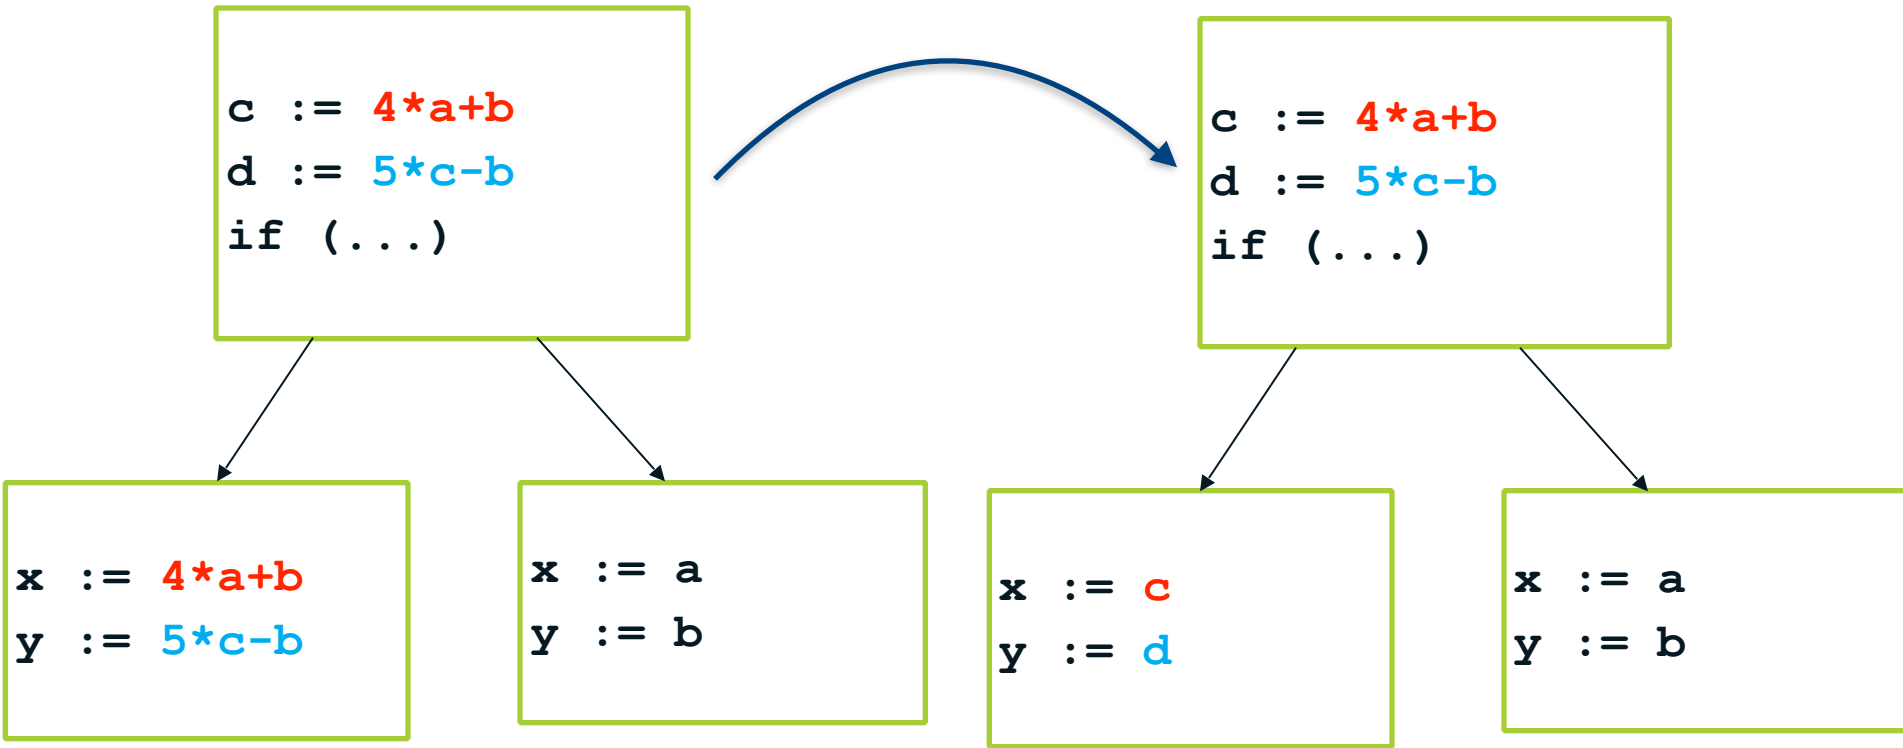
Mid End Optimizer – Clean up



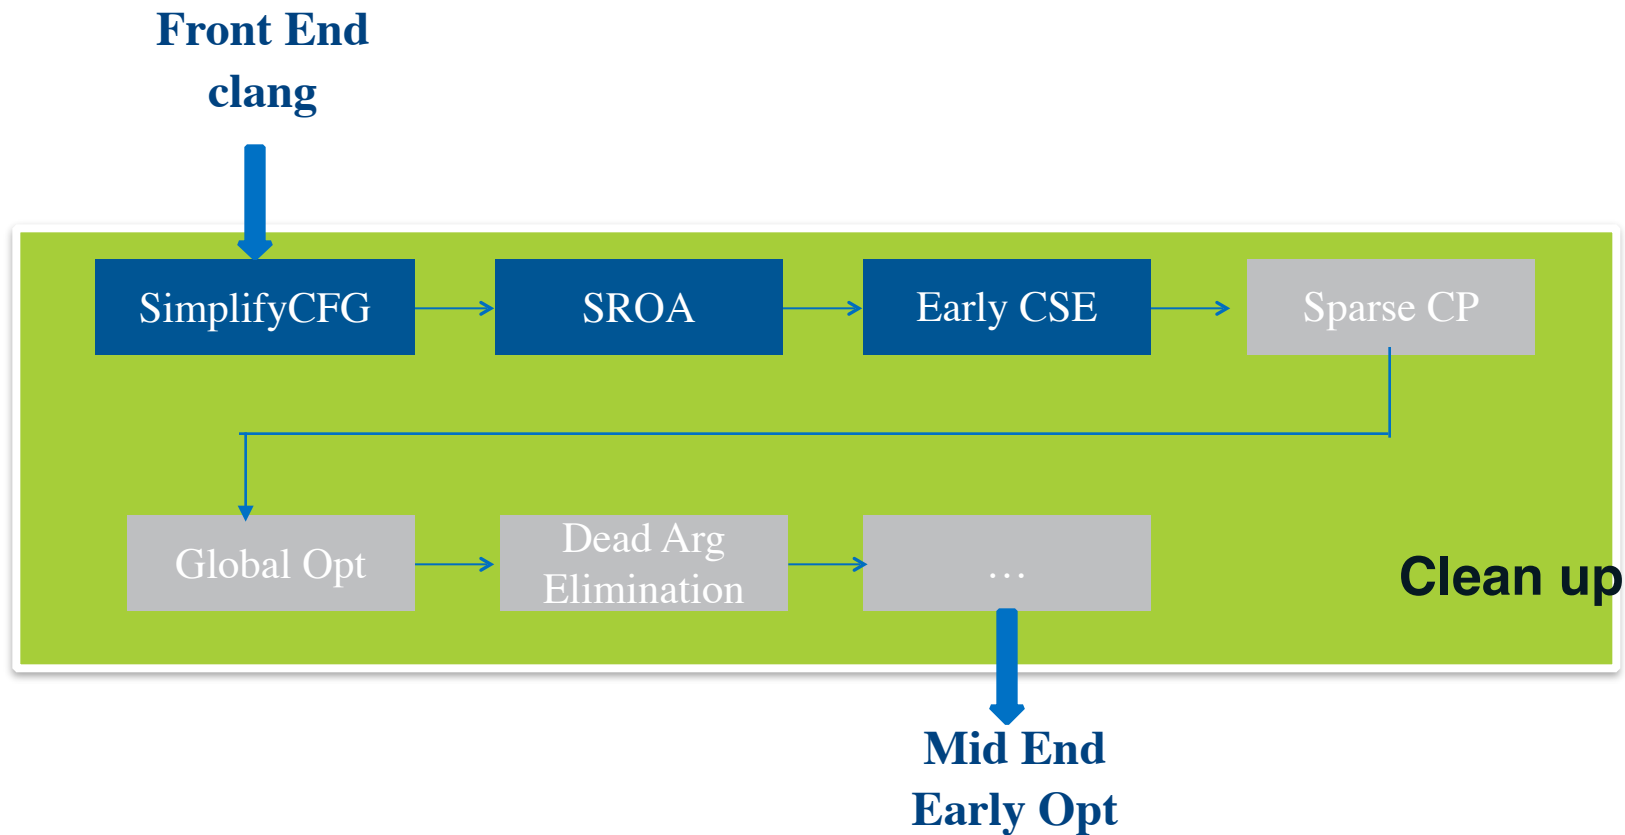
Mid End Optimizer – Clean up



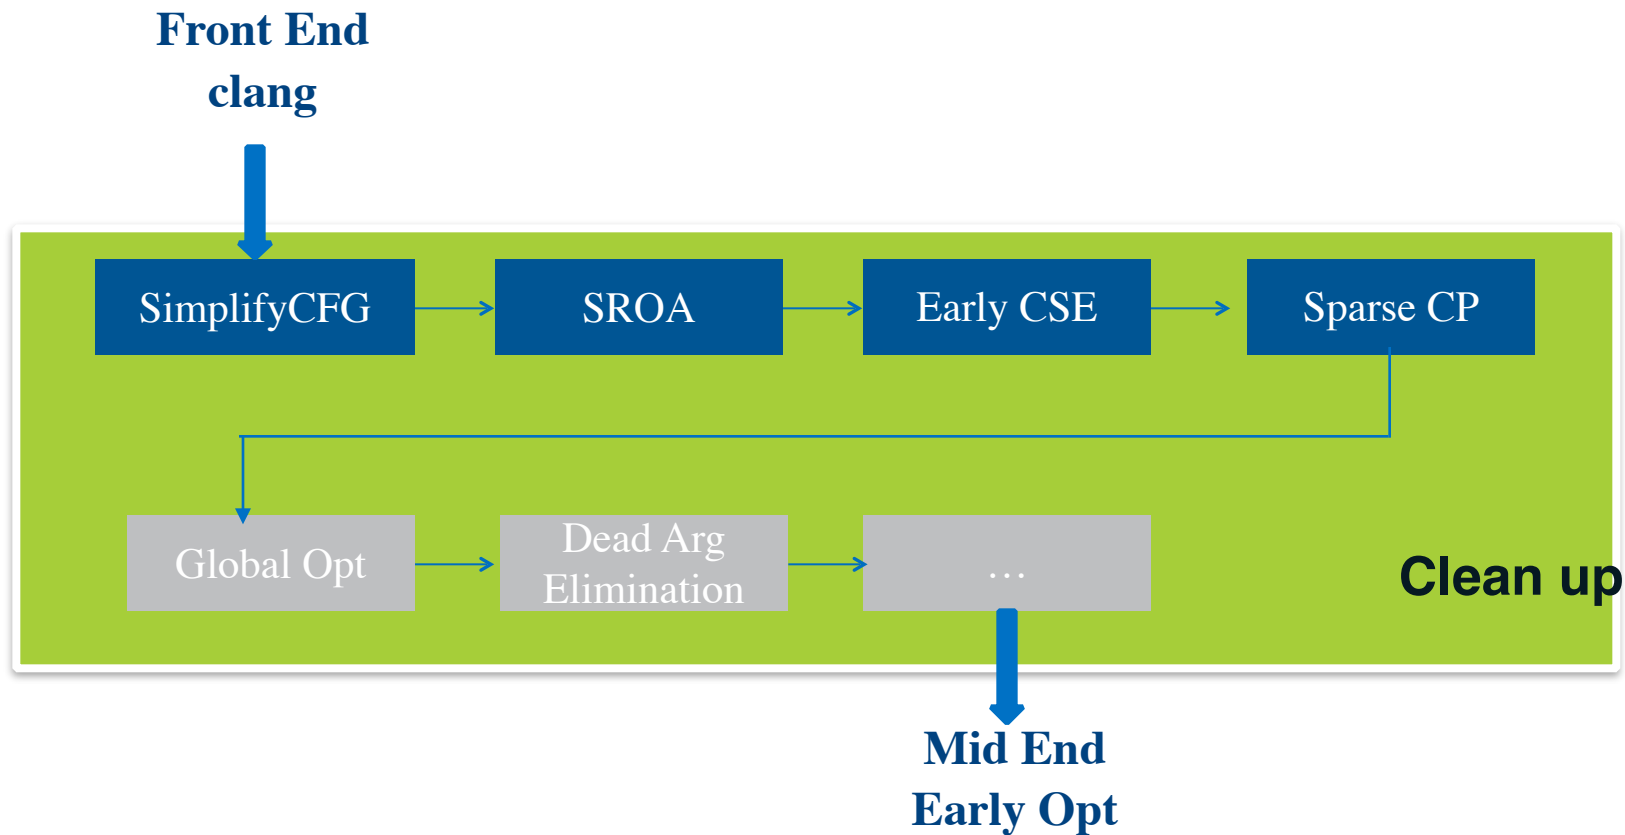
Common Subexpression Elimination



Mid End Optimizer – Clean up



Mid End Optimizer – Clean up



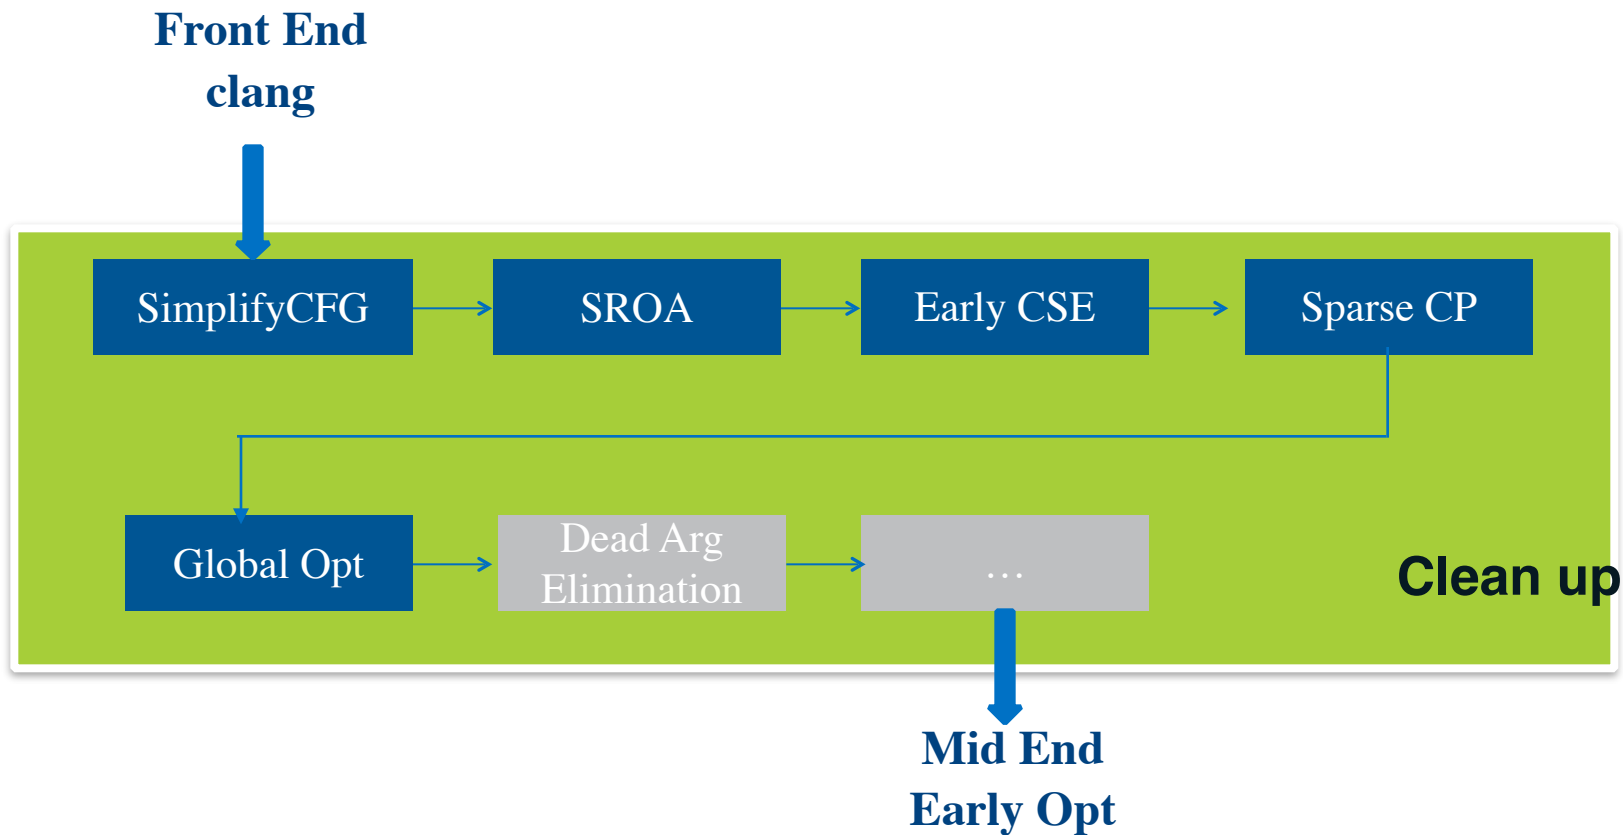
Constant Propagation

```
x := 5  
z := y + x;
```

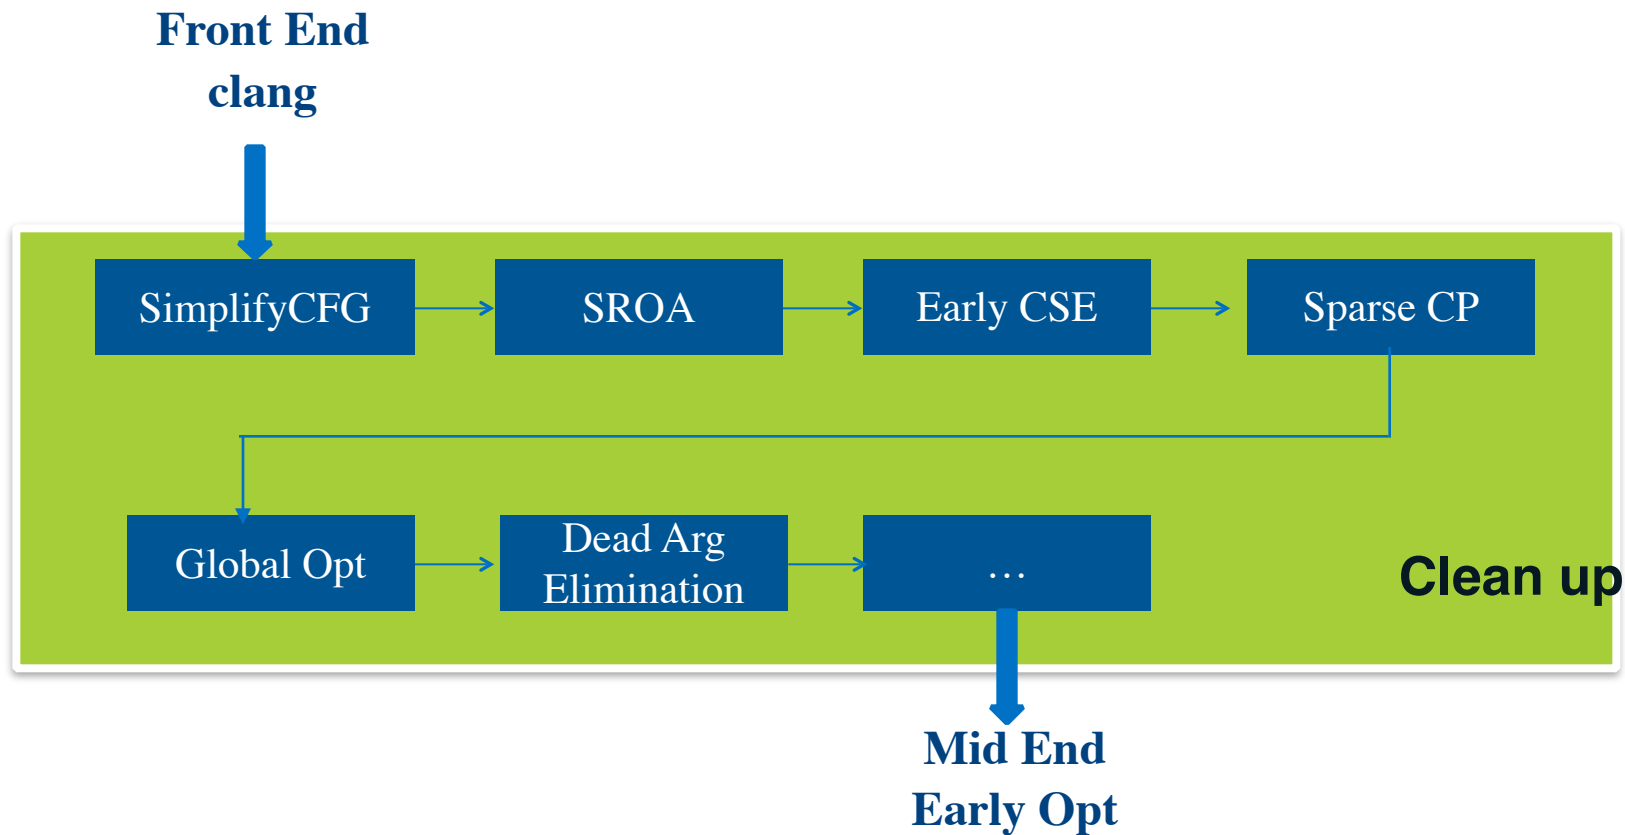


```
x := 5  
z := y + 5;
```

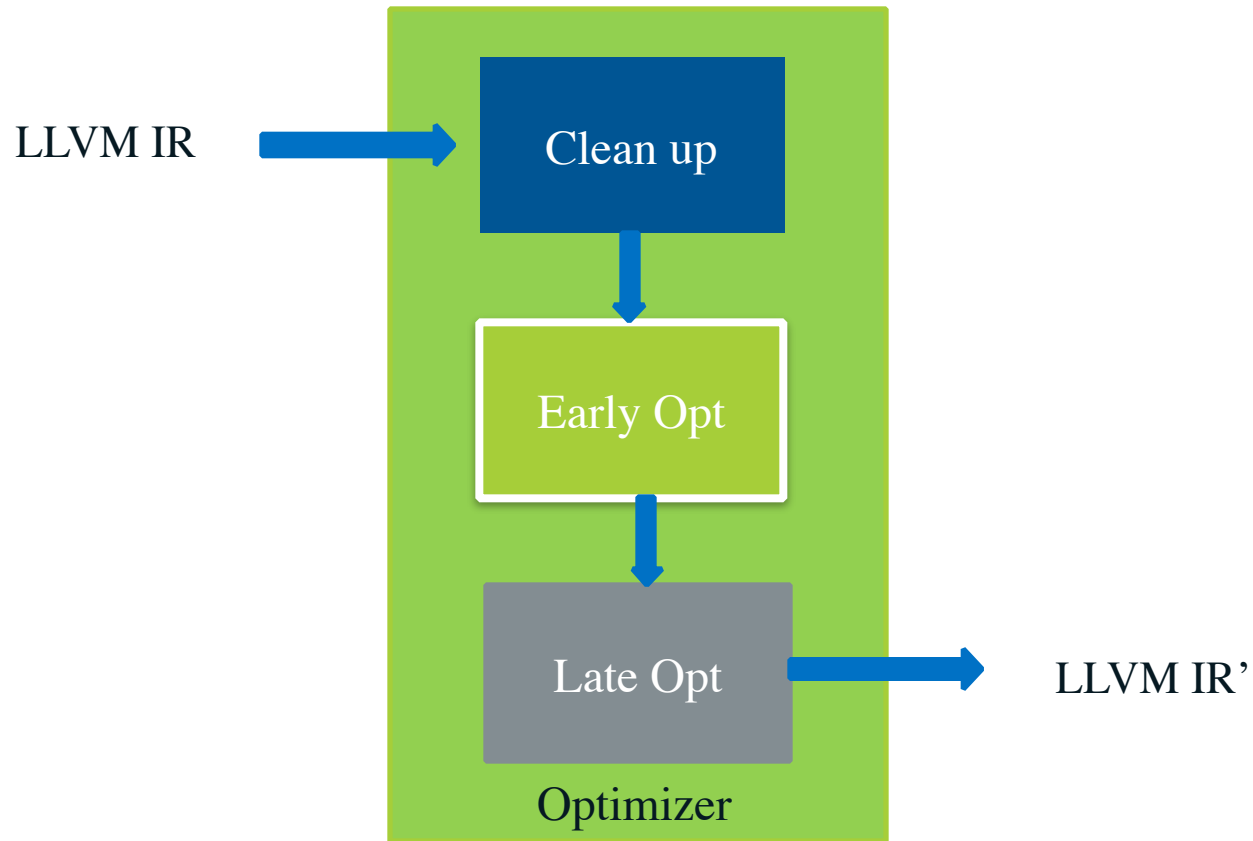

Mid End Optimizer – Clean up



Mid End Optimizer – Clean up



Mid End Optimizer – Early Opt



Mid End Optimizer – Early Opt

Clean up

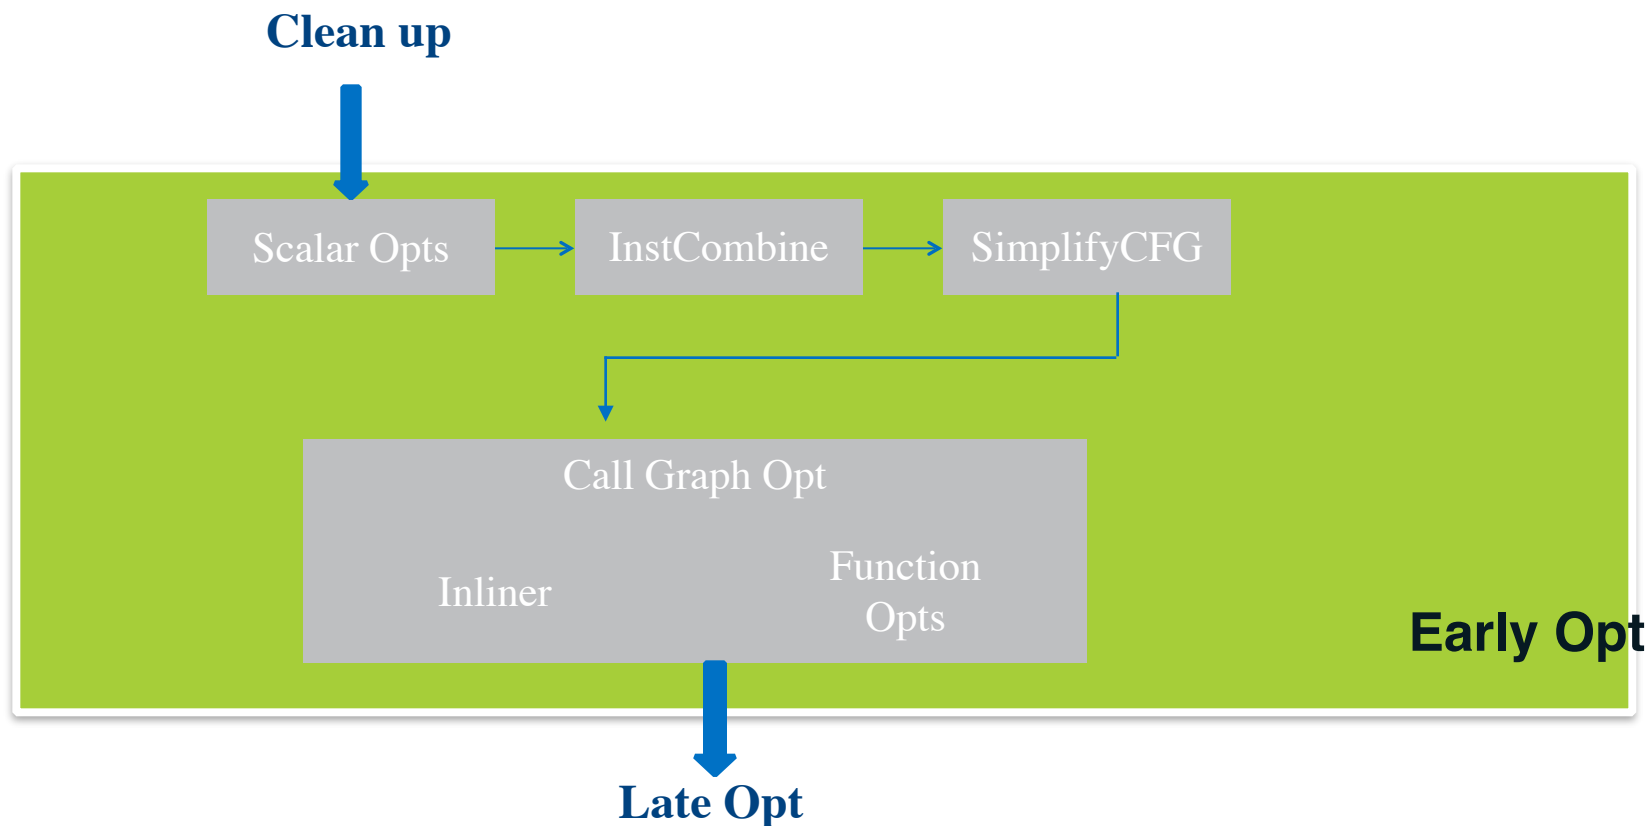


Early Opt

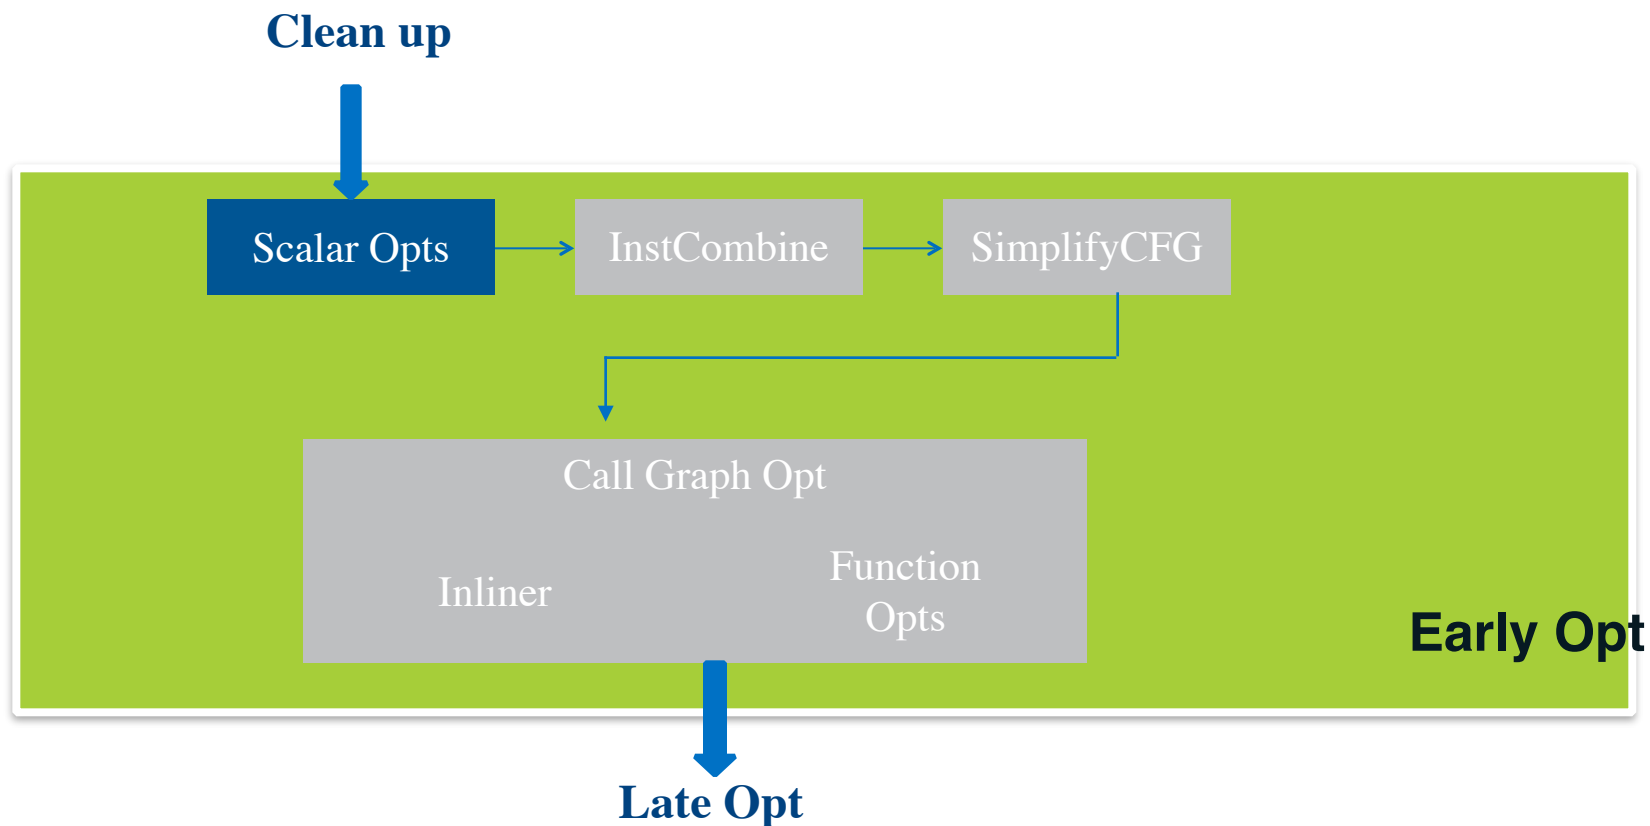


Late Opt

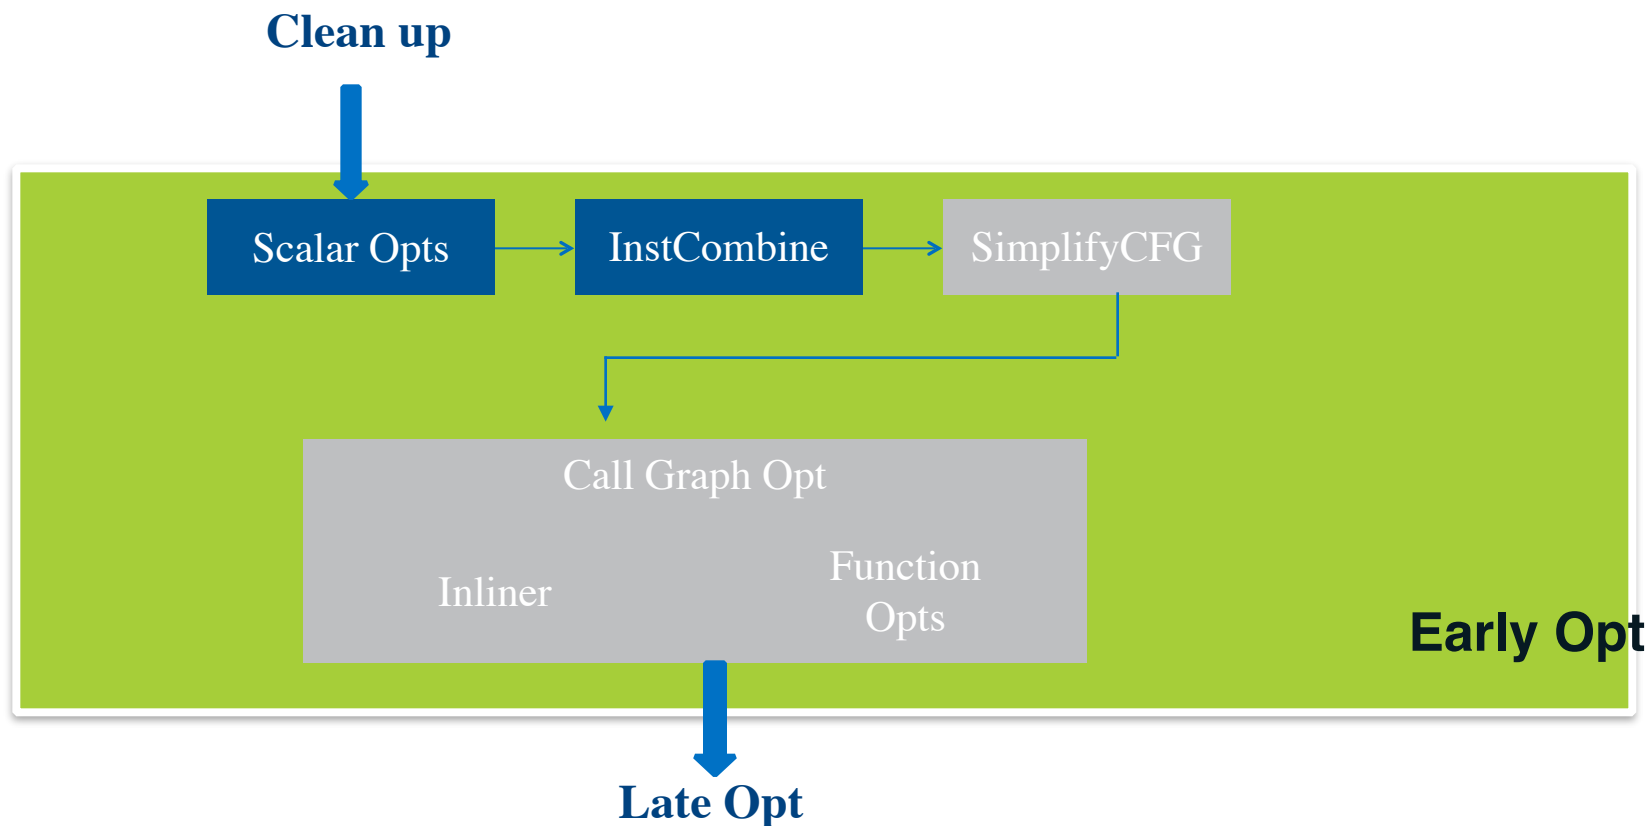
Mid End Optimizer – Early Opt



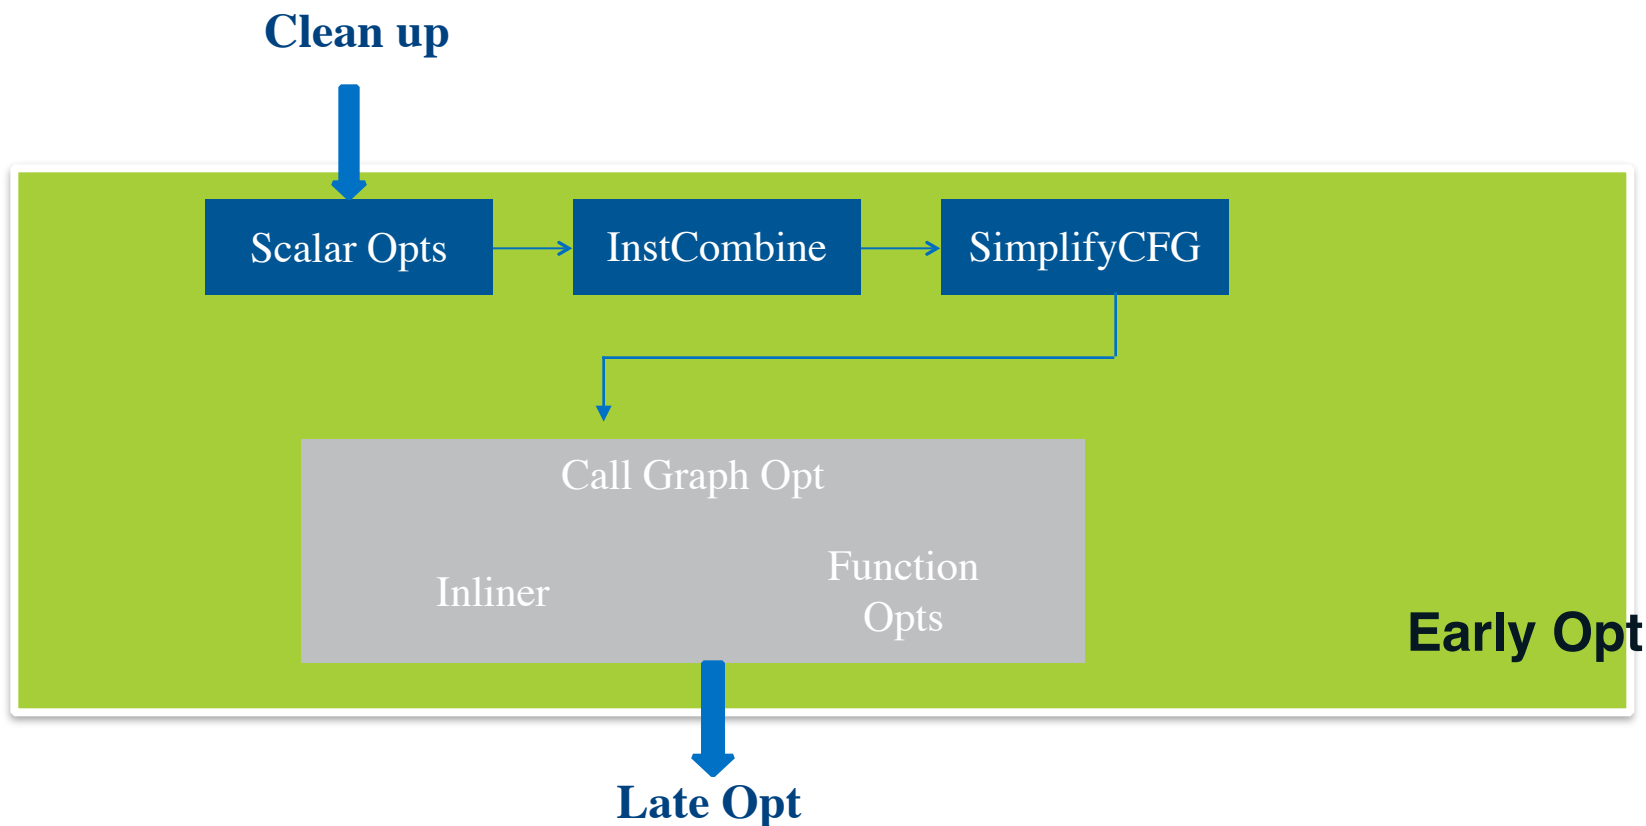
Mid End Optimizer – Early Opt



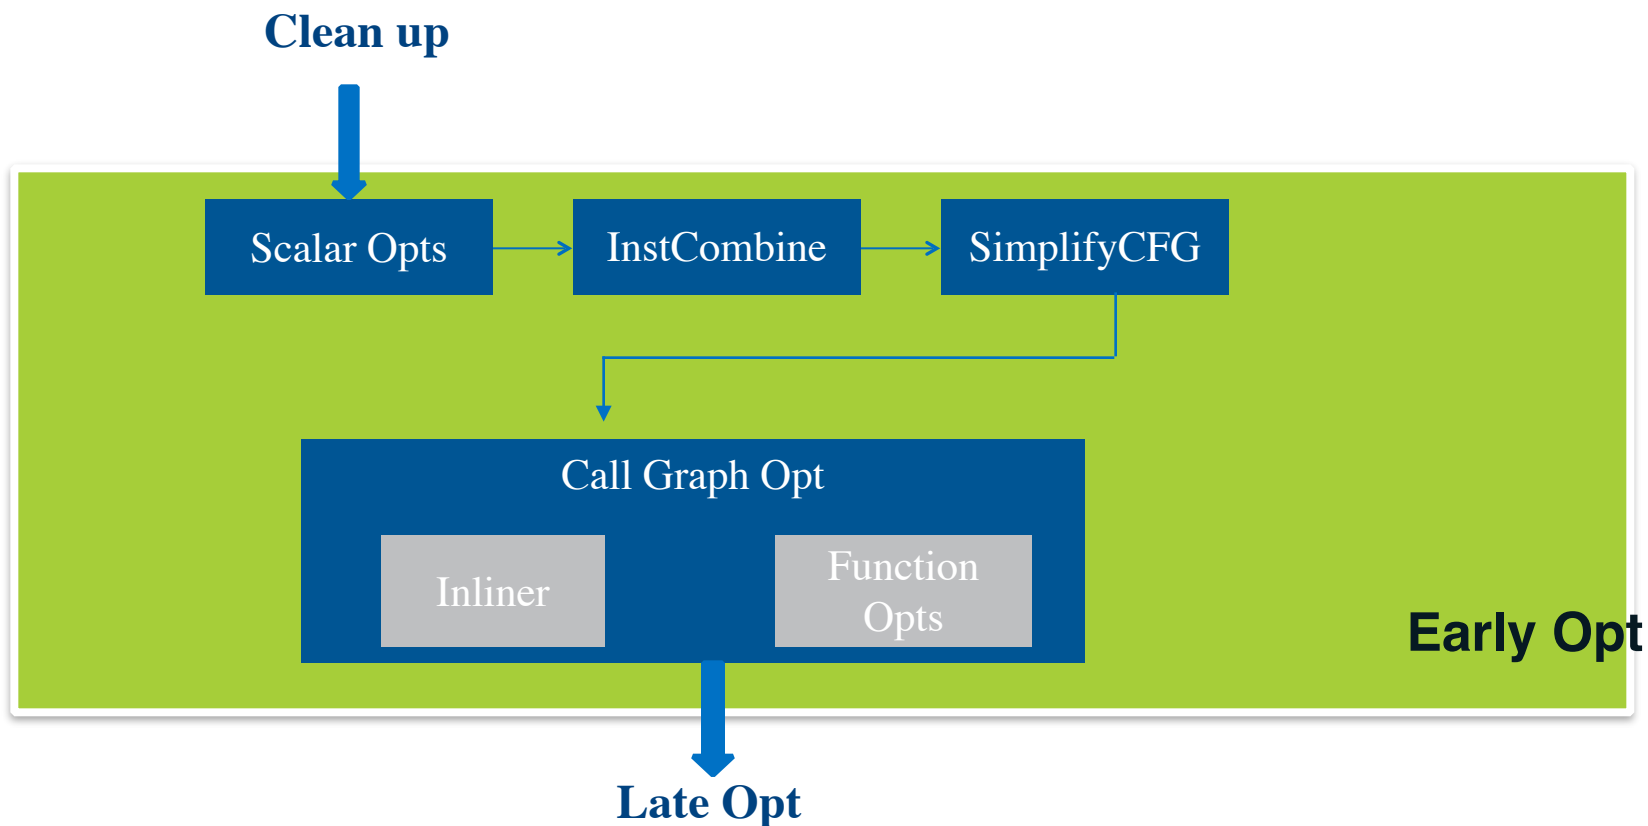
Mid End Optimizer – Early Opt



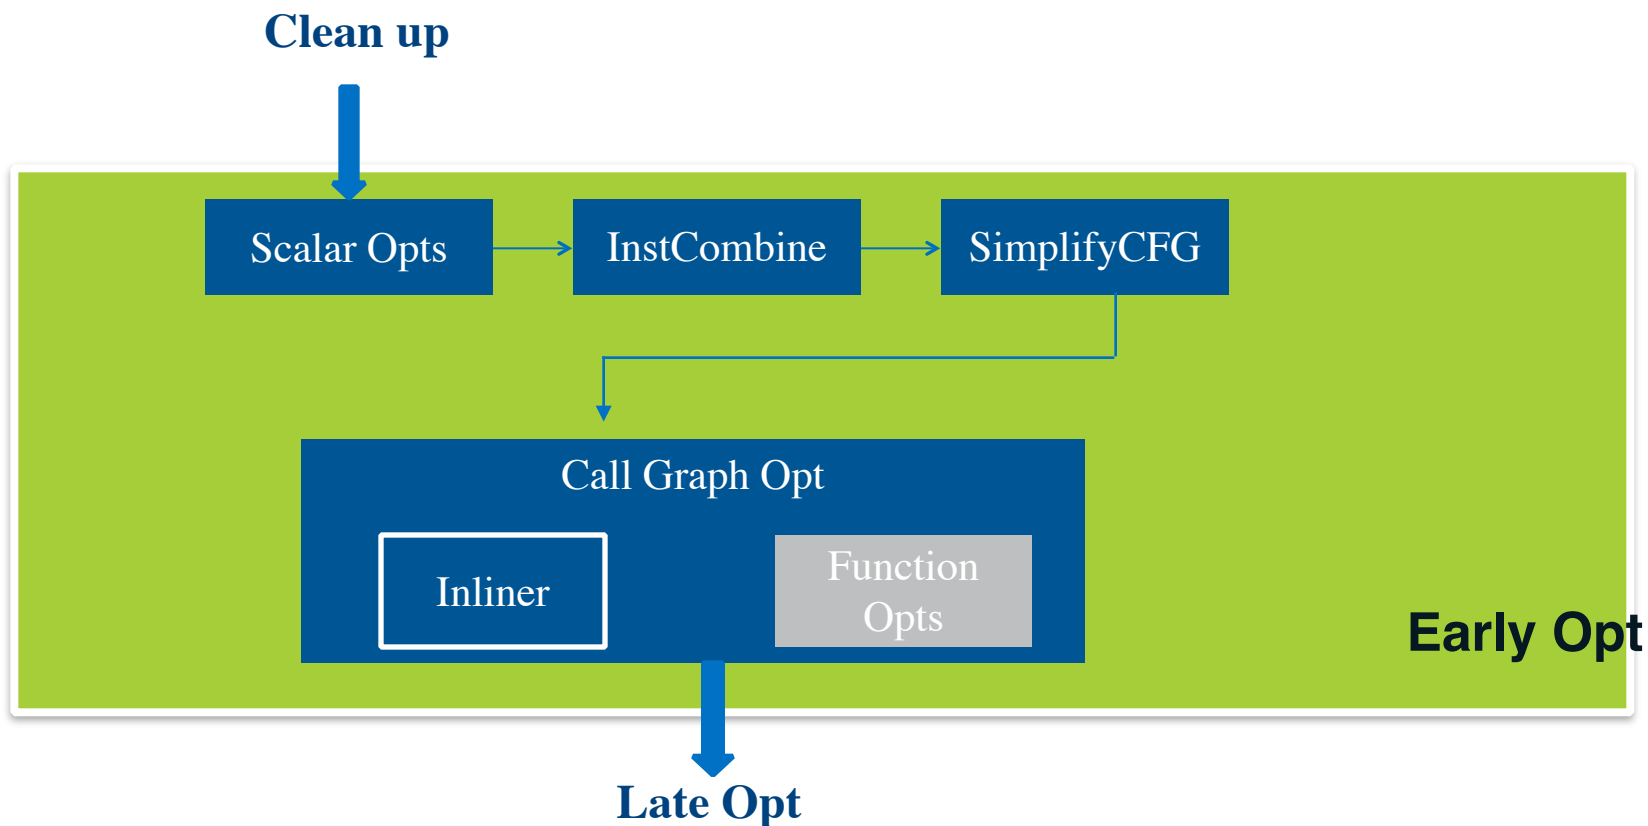
Mid End Optimizer – Early Opt



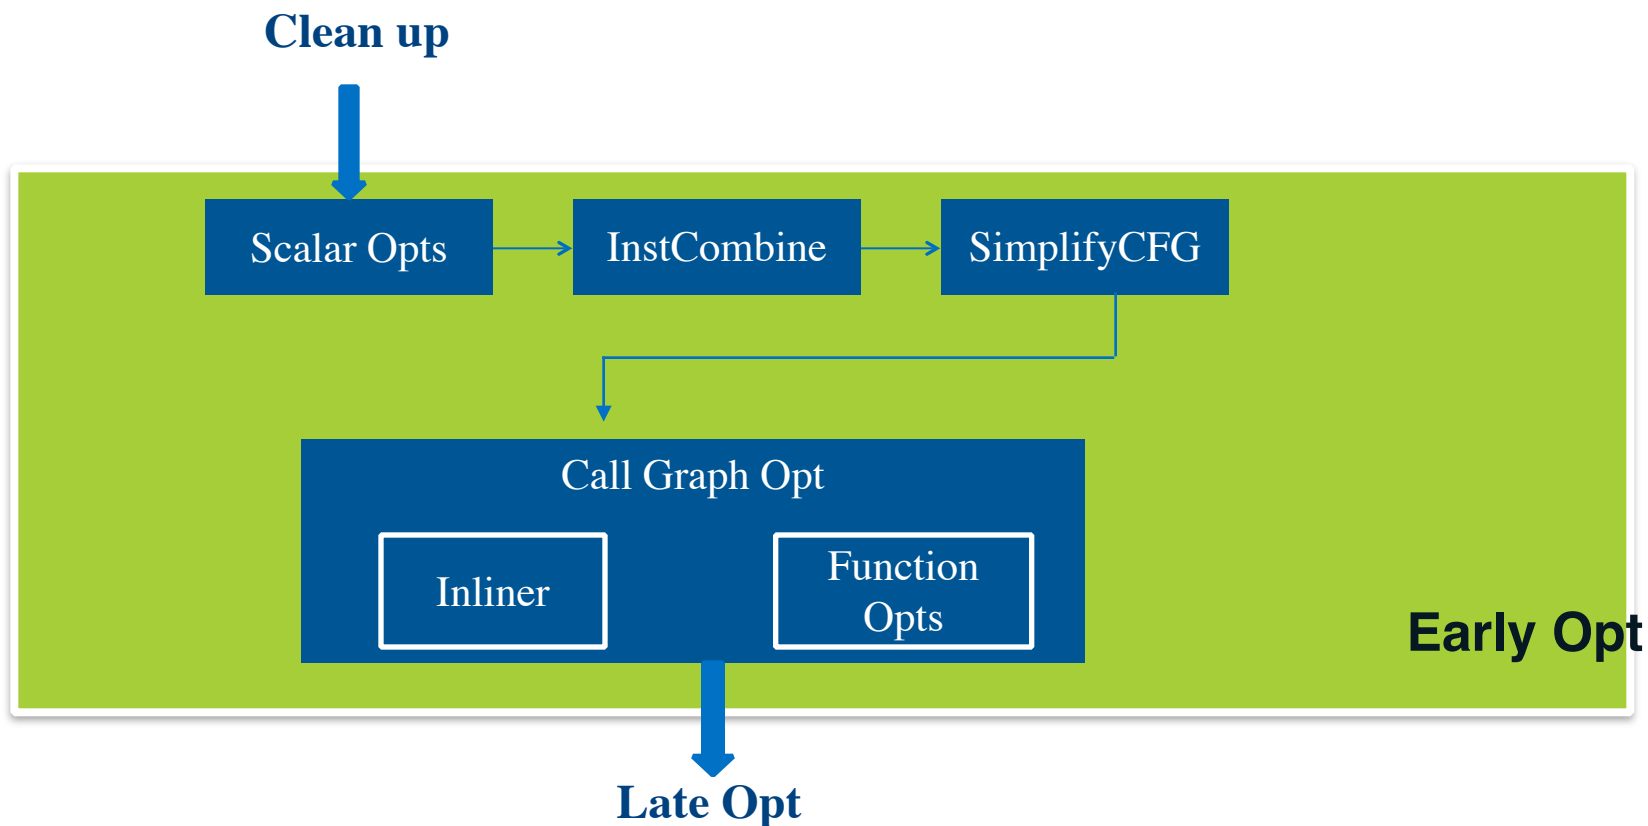
Mid End Optimizer – Early Opt



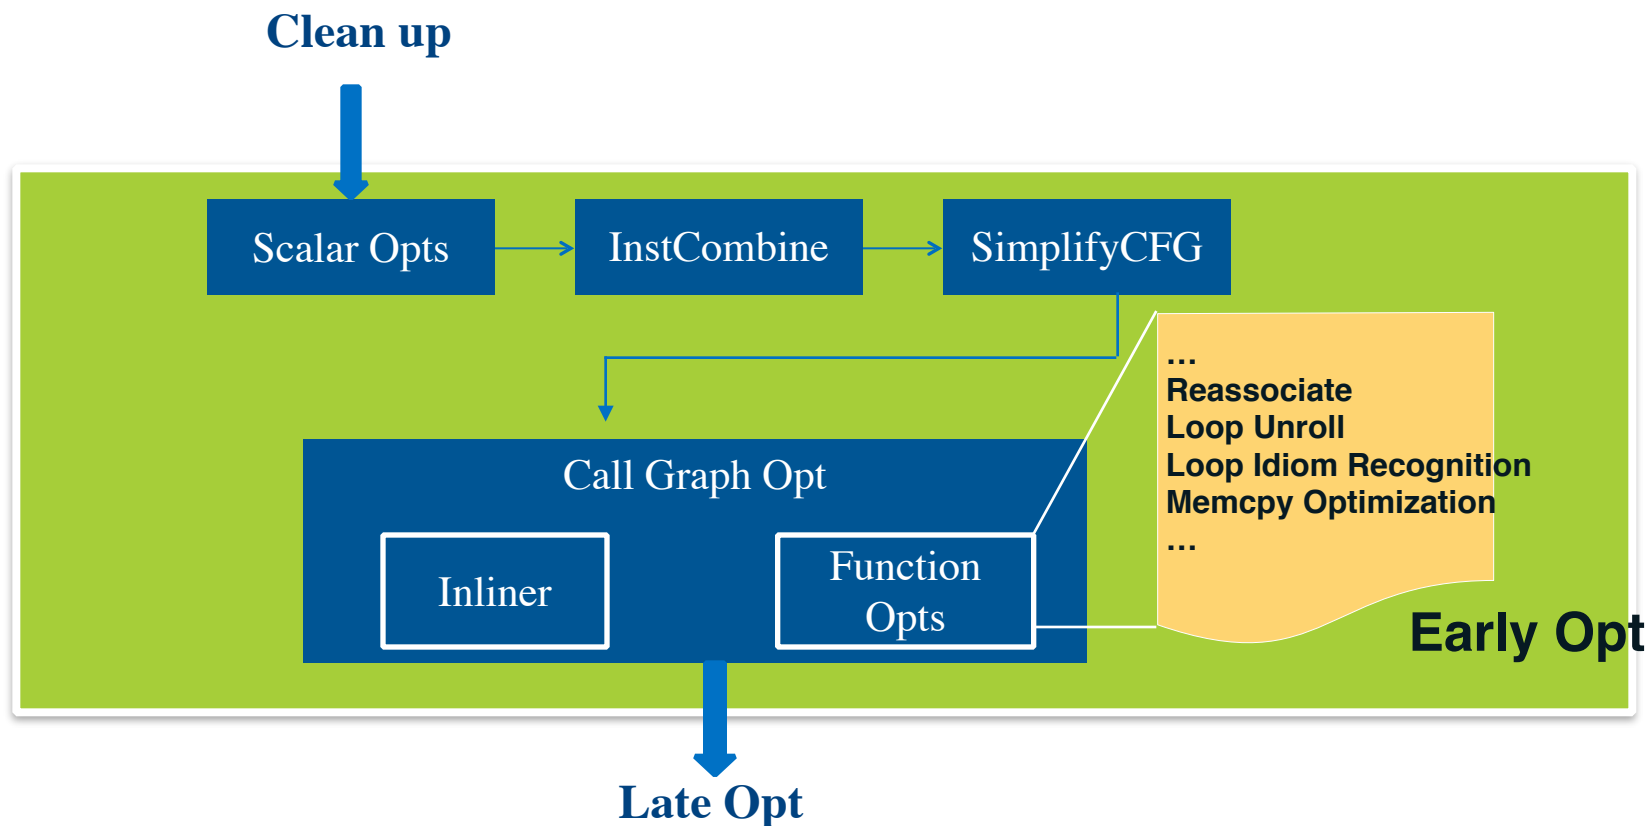
Mid End Optimizer – Early Opt



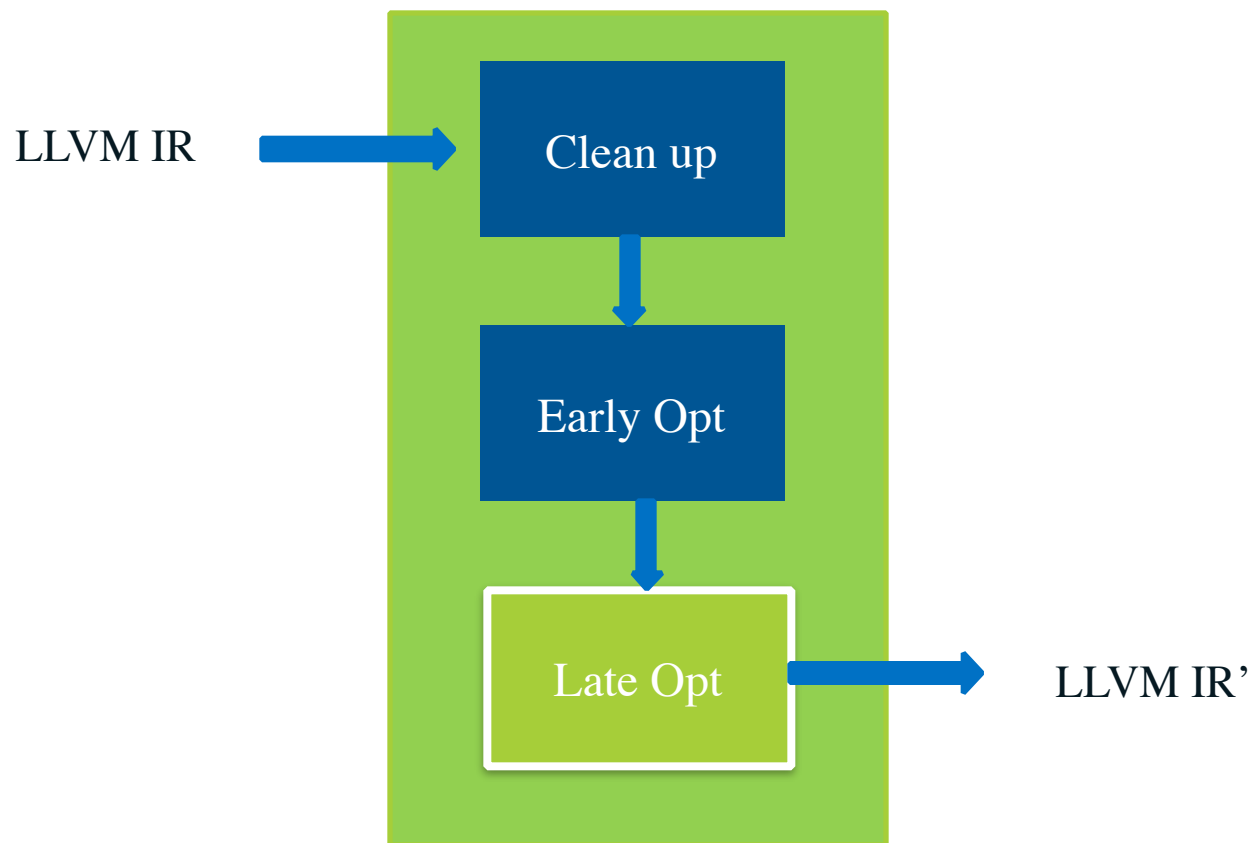
Mid End Optimizer – Early Opt




Mid End Optimizer – Early Opt



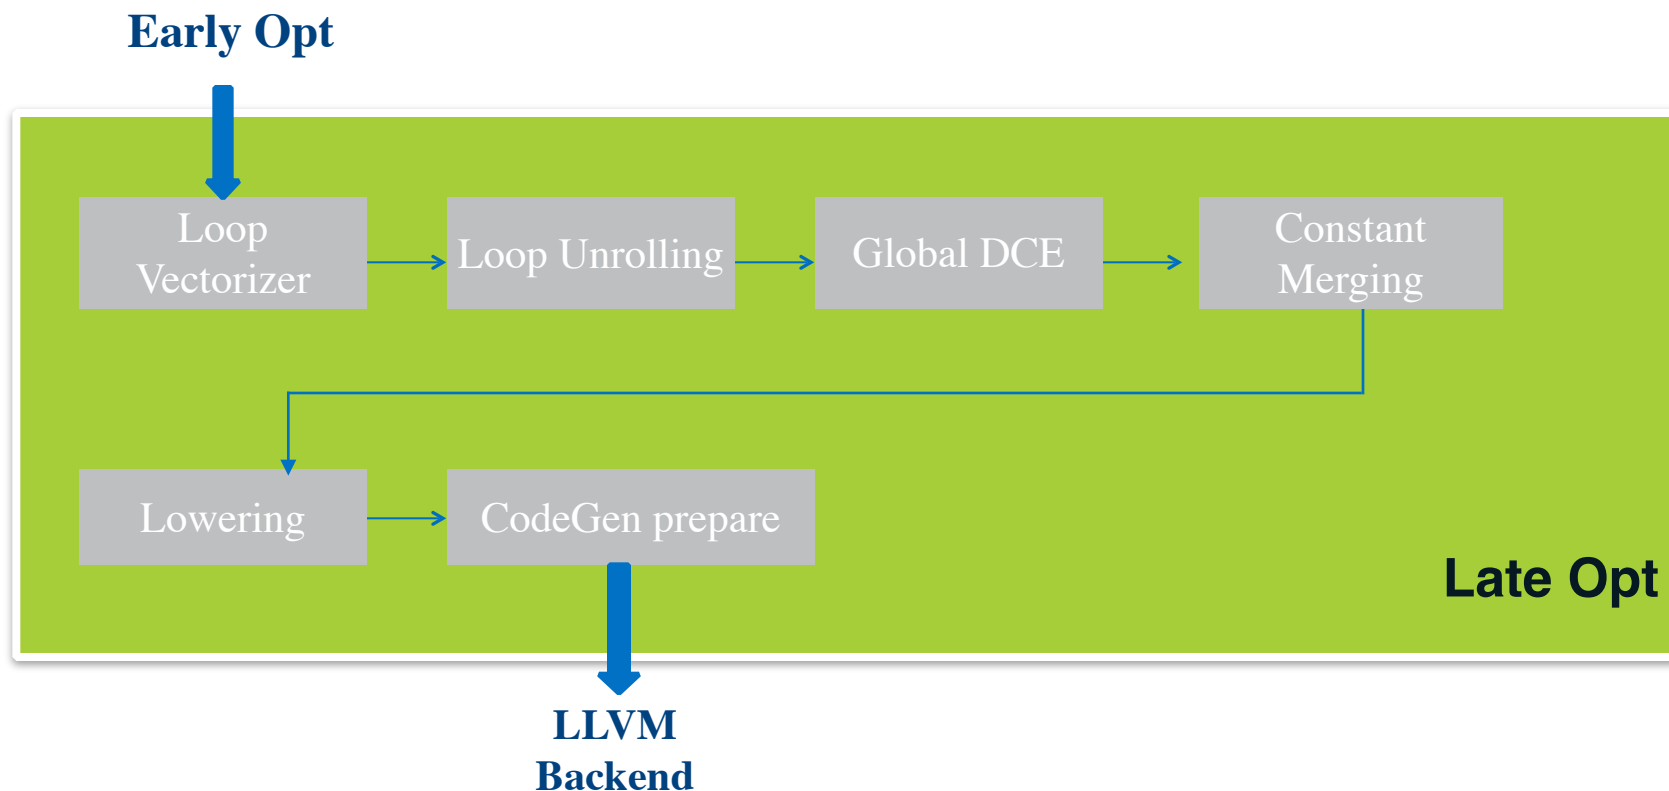
Mid End Optimizer – Late Opt



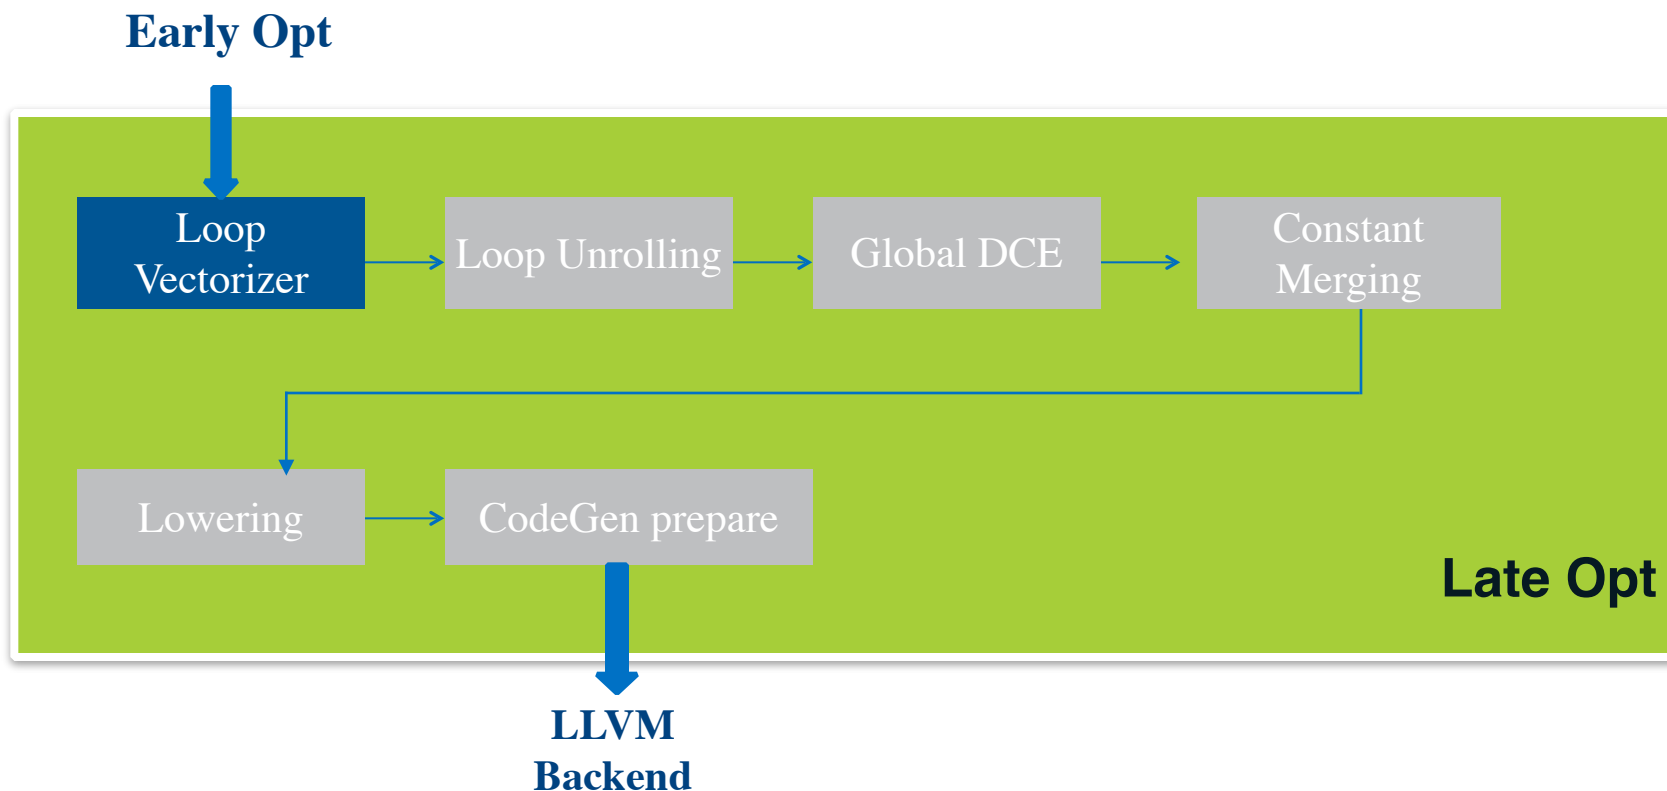
Mid End Optimizer – Late Opt



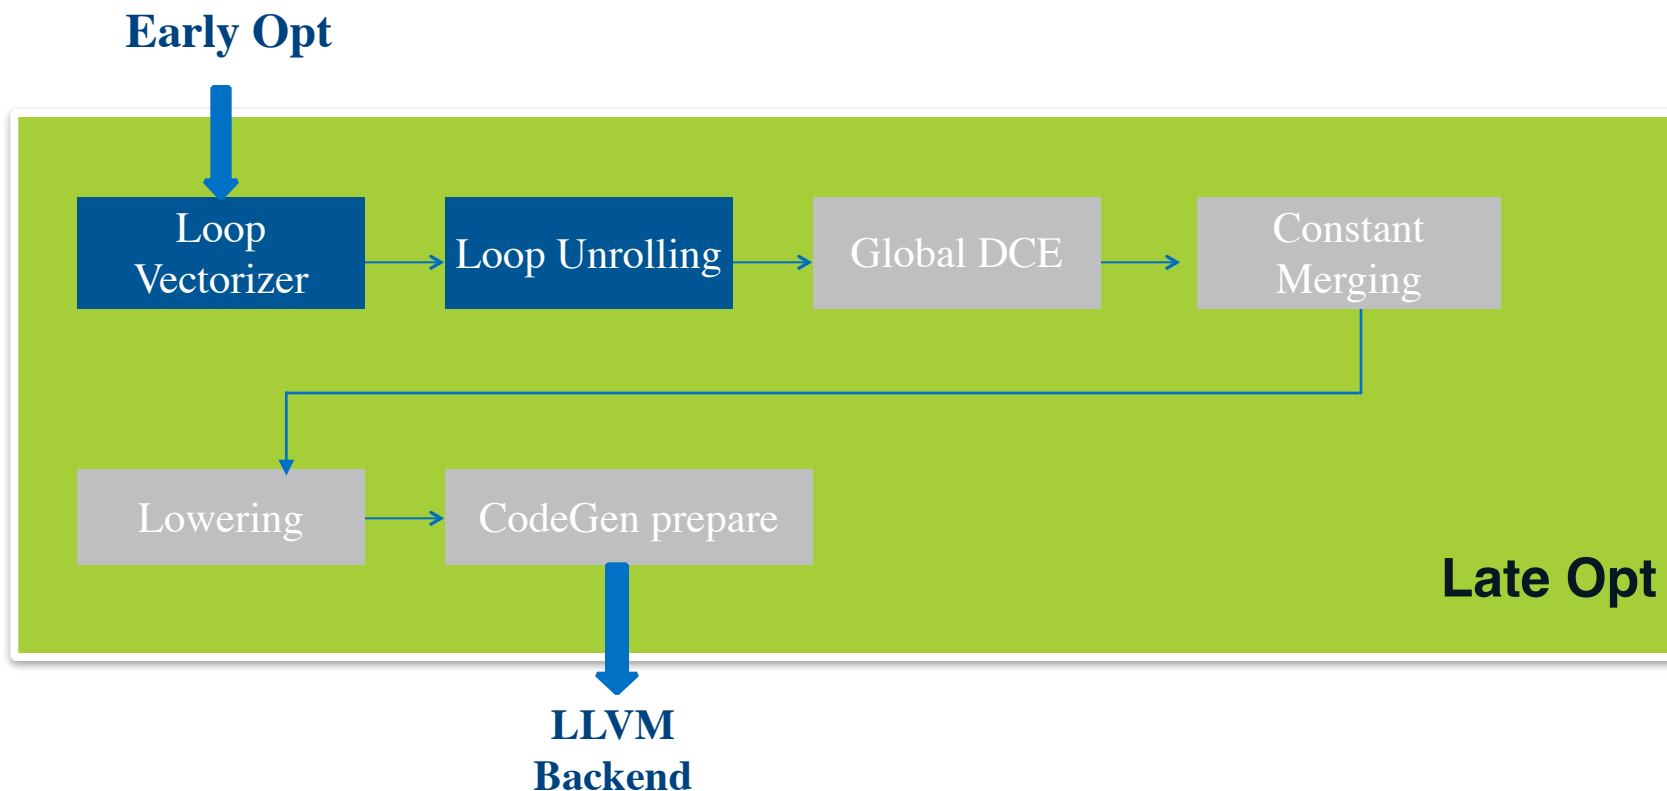
Mid End Optimizer – Late Opt



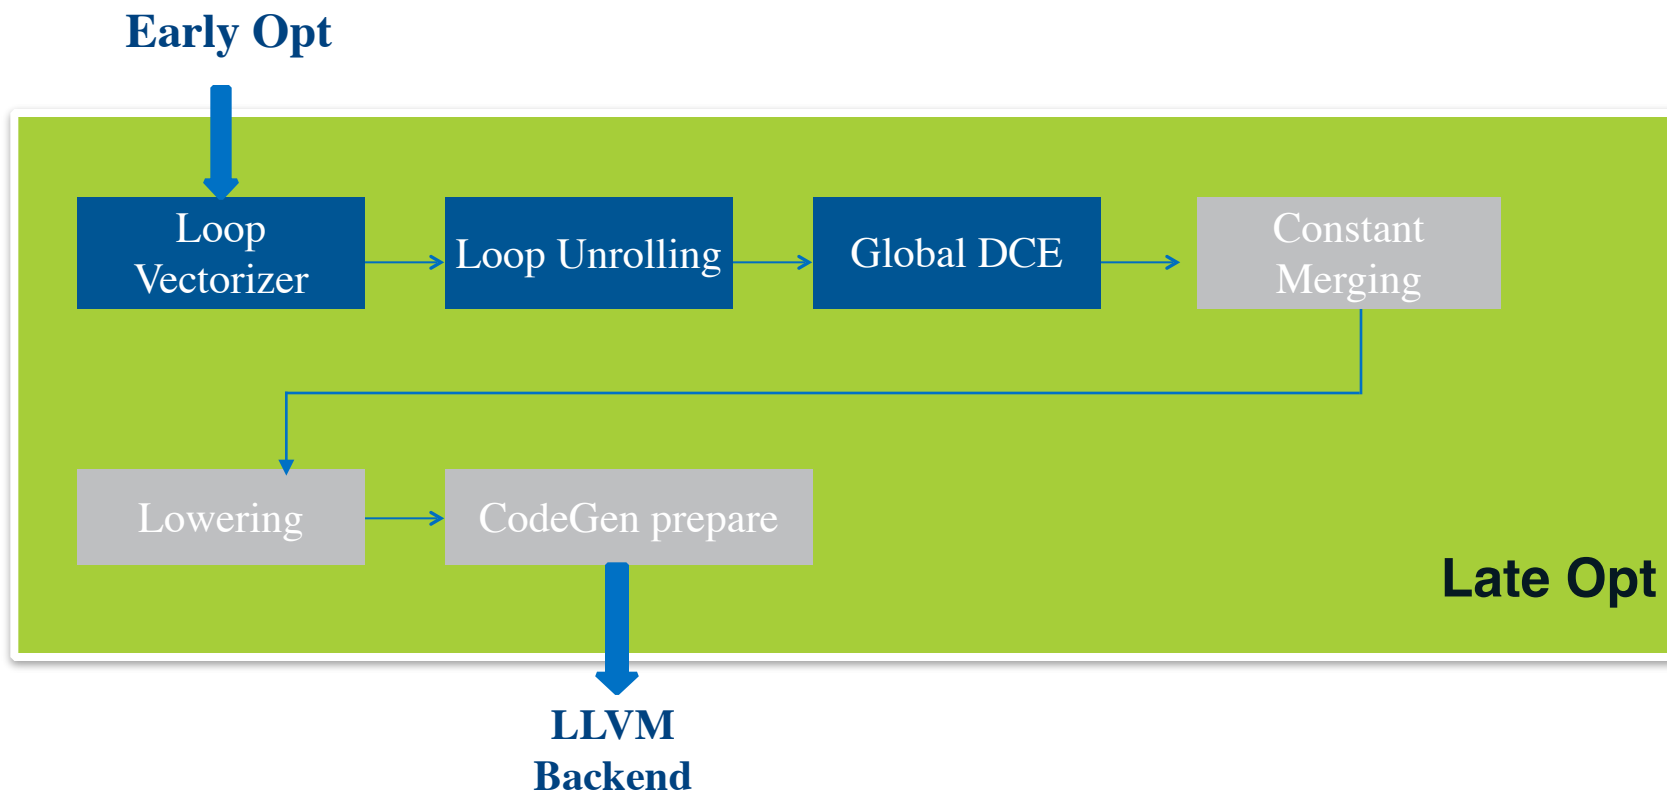
Mid End Optimizer – Late Opt



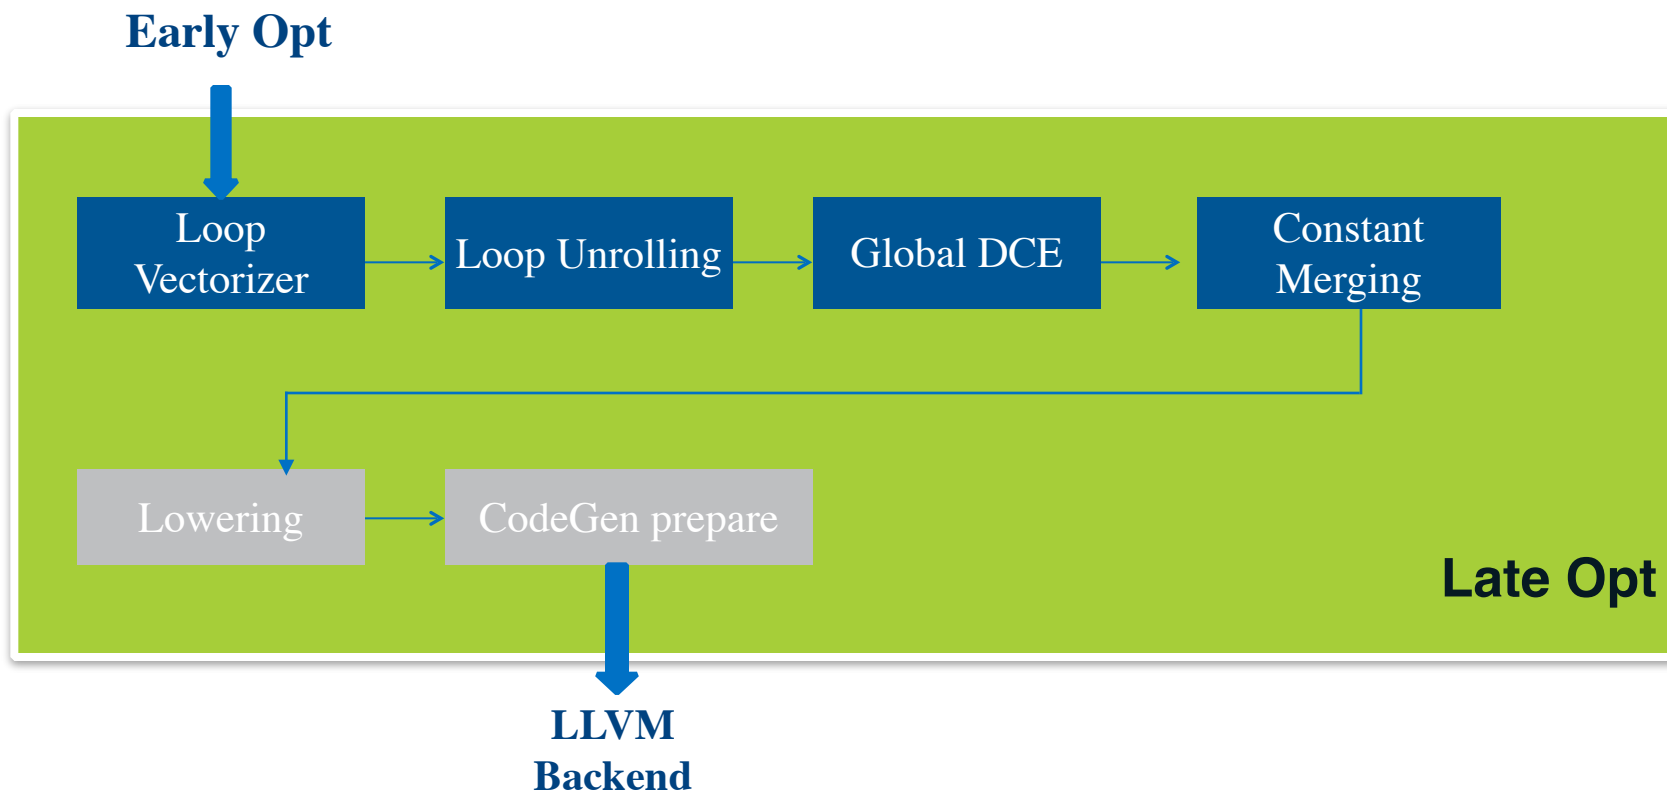
Mid End Optimizer – Late Opt



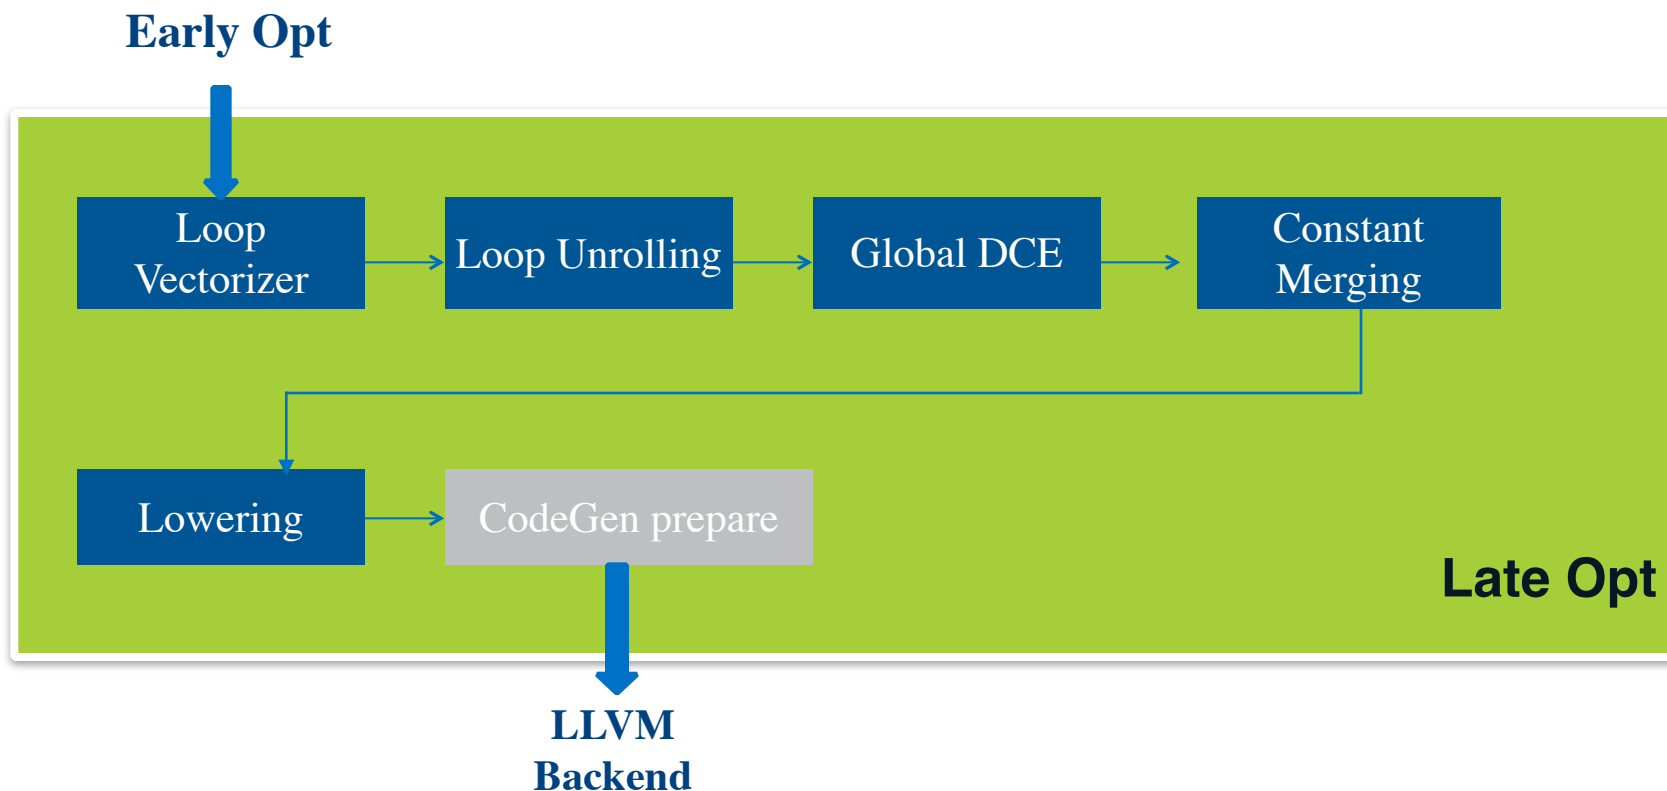
Mid End Optimizer – Late Opt



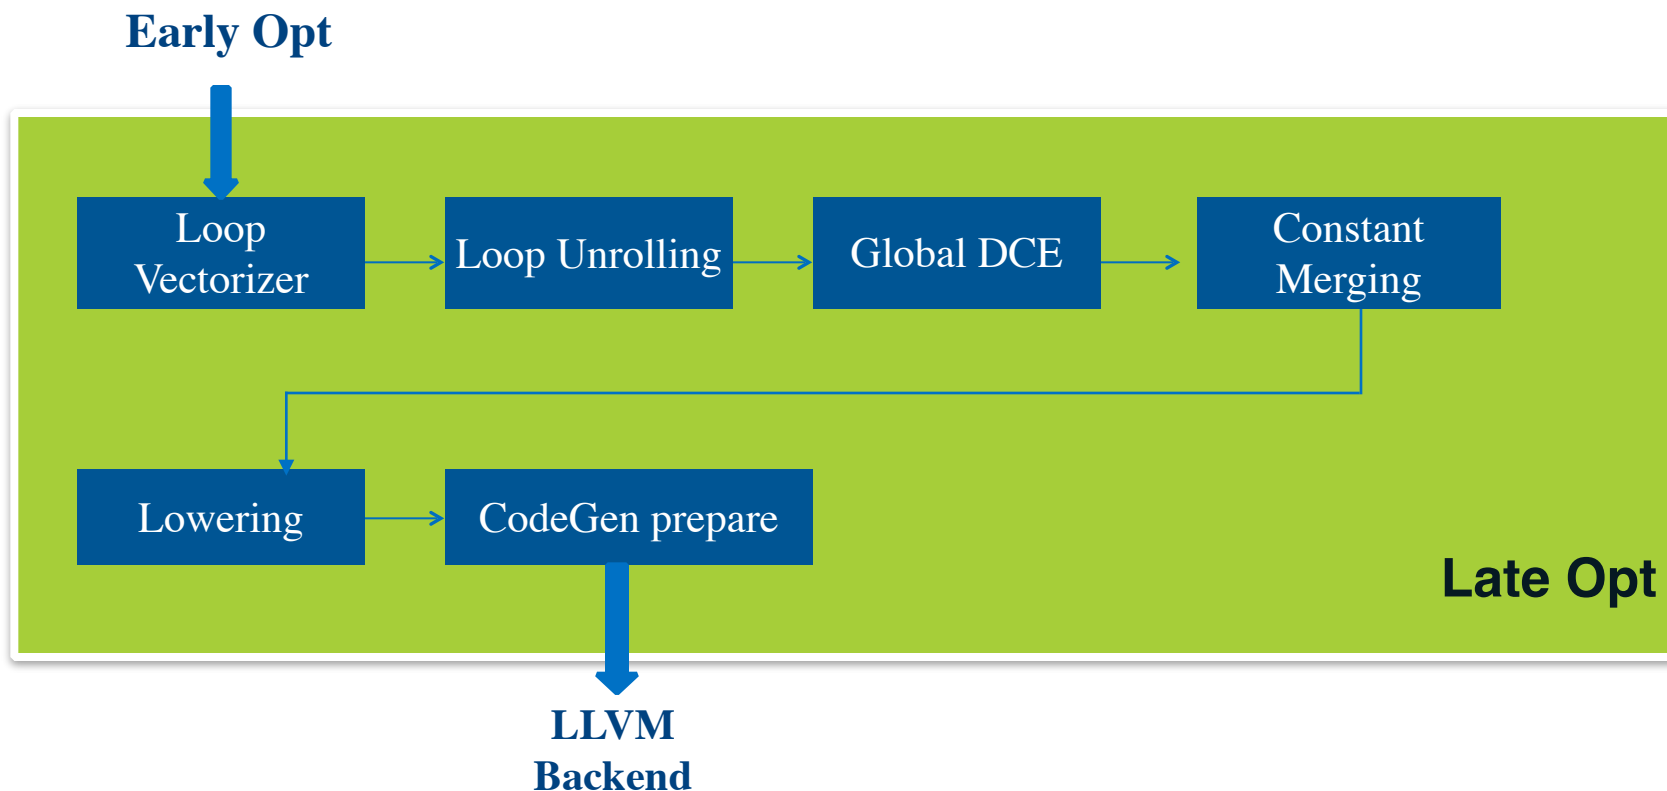
Mid End Optimizer – Late Opt




Mid End Optimizer – Late Opt



Mid End Optimizer – Late Opt




LLVM Pipeline



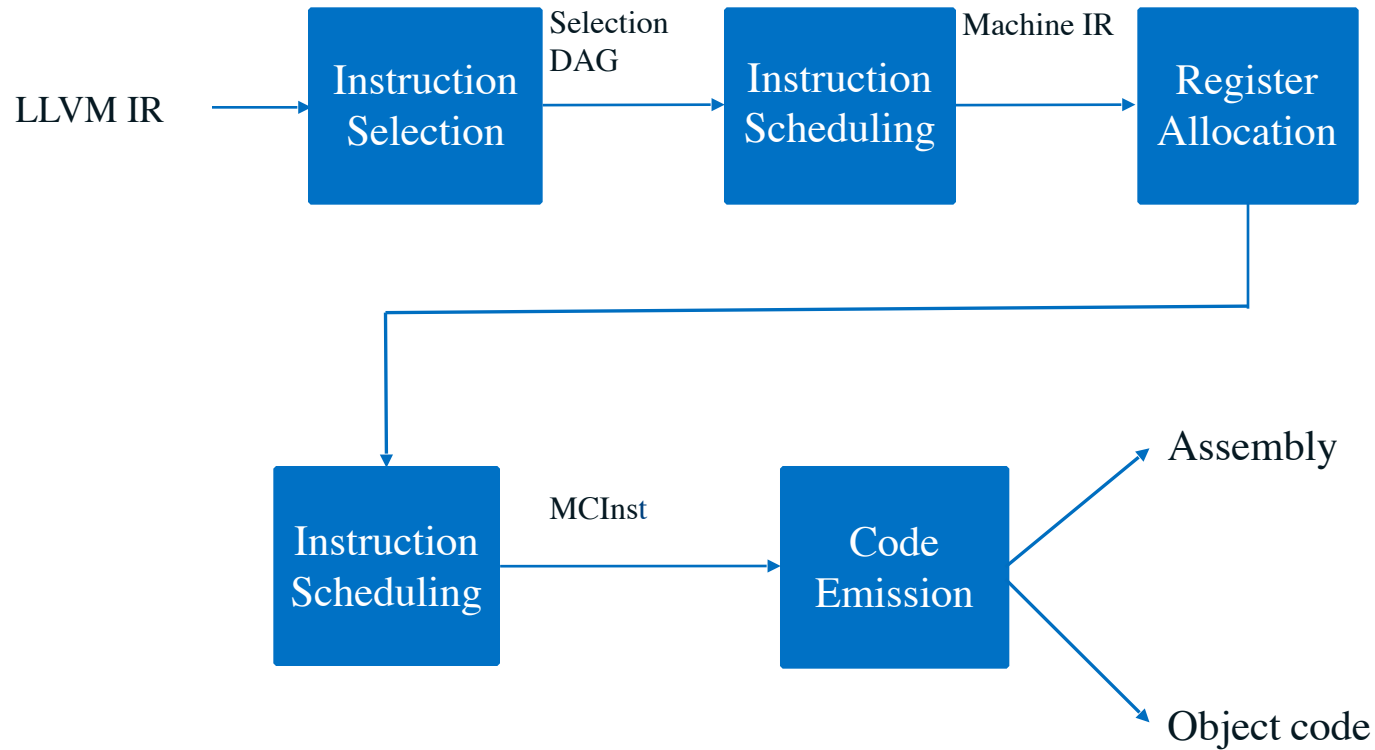
LLVM Backend

LLVM Backend

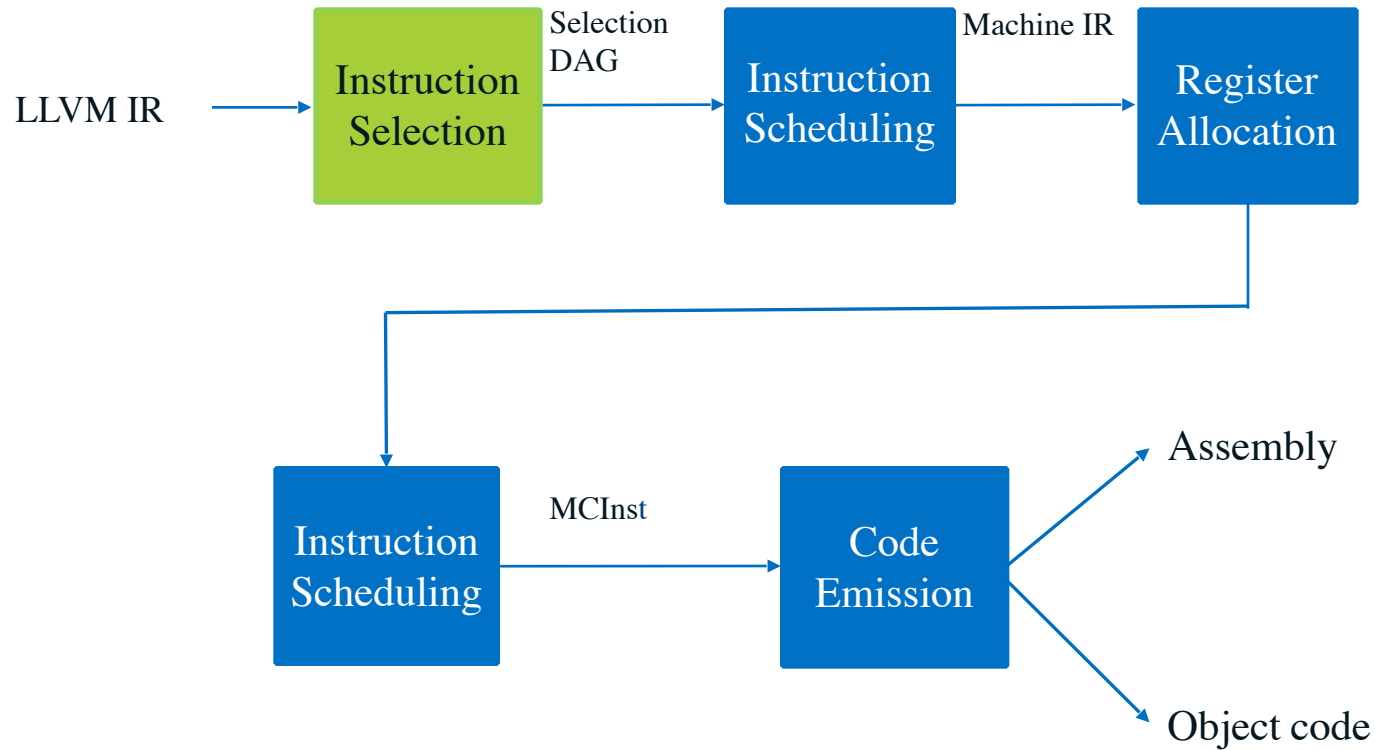
- Code generation for specific architectures via
 - Instruction Selection
 - Instruction Scheduling
 - Register Allocation



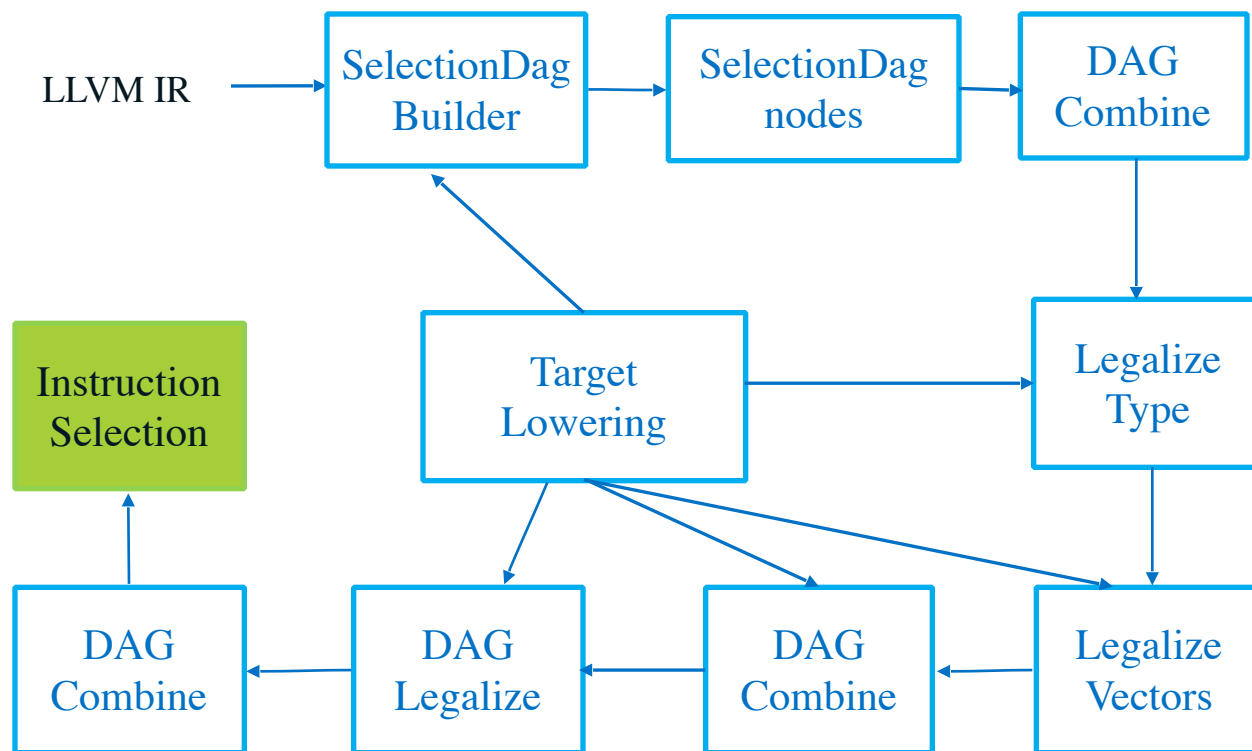
LLVM Backend Pipeline



LLVM Backend Pipeline



LLVM Backend - Instruction Selection Pipeline



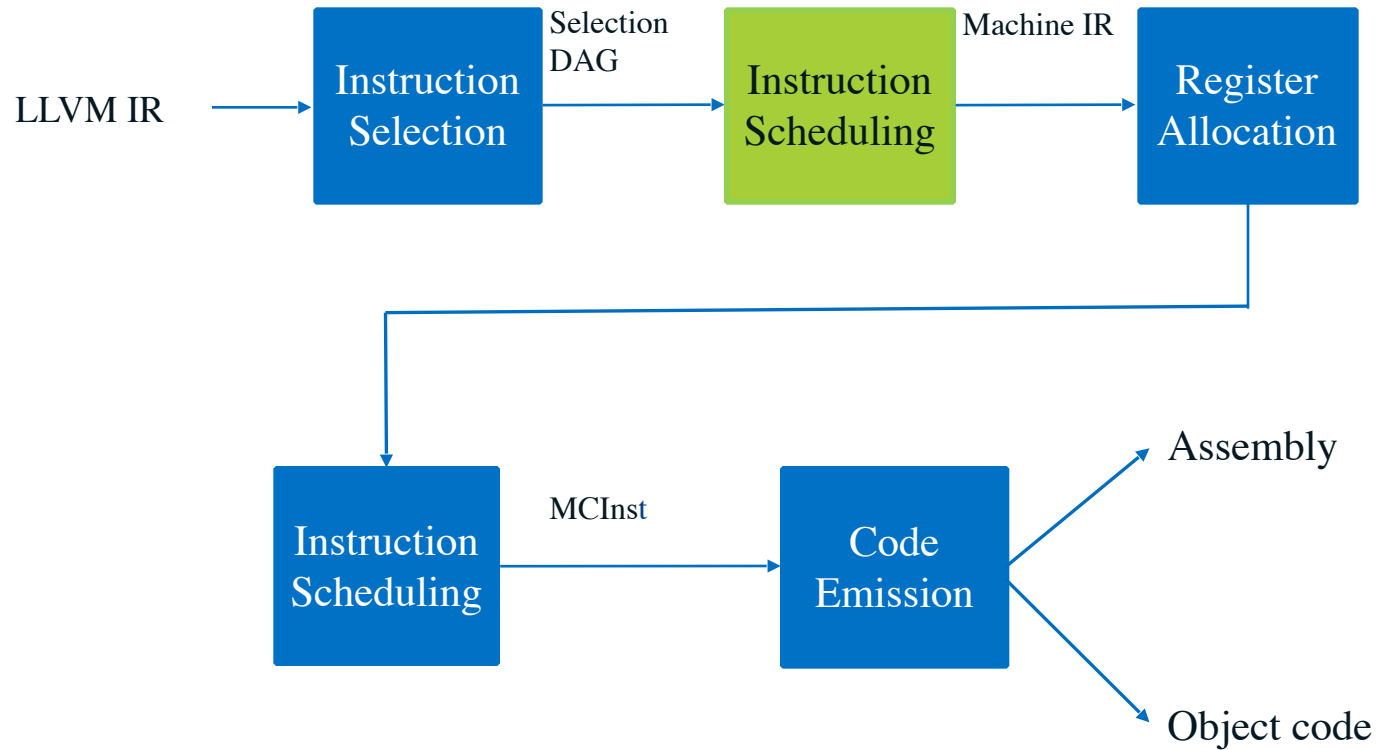
Selection DAG Instruction Select Phase Example

```
%t1 = add i32 %W, %X  
%t2 = mul i32 %t1, %Y  
%t3 = add i32 %t2, %Z
```

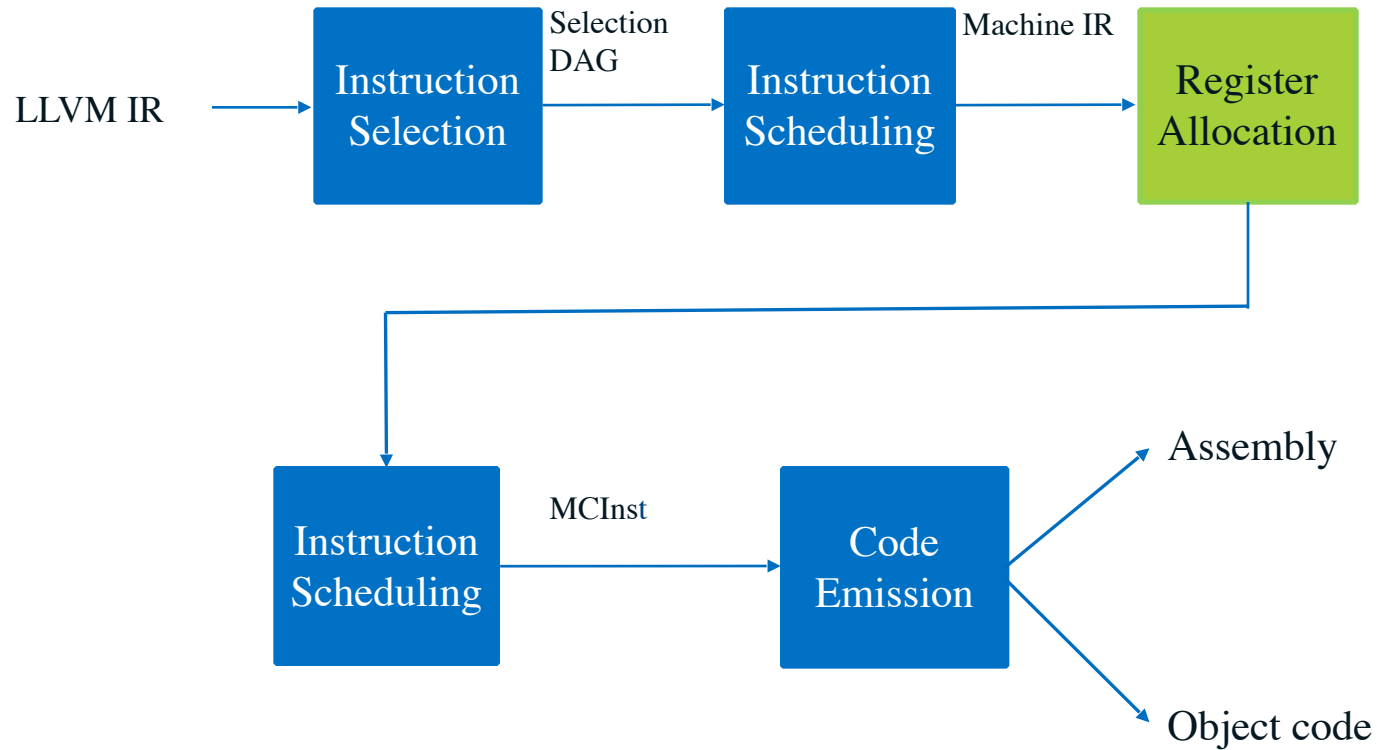
```
(add:i32 (mul:i32 (add:i32 W, X), Y), Z)
```

```
(MADDrr (ADDrr W, X), Y, Z)
```

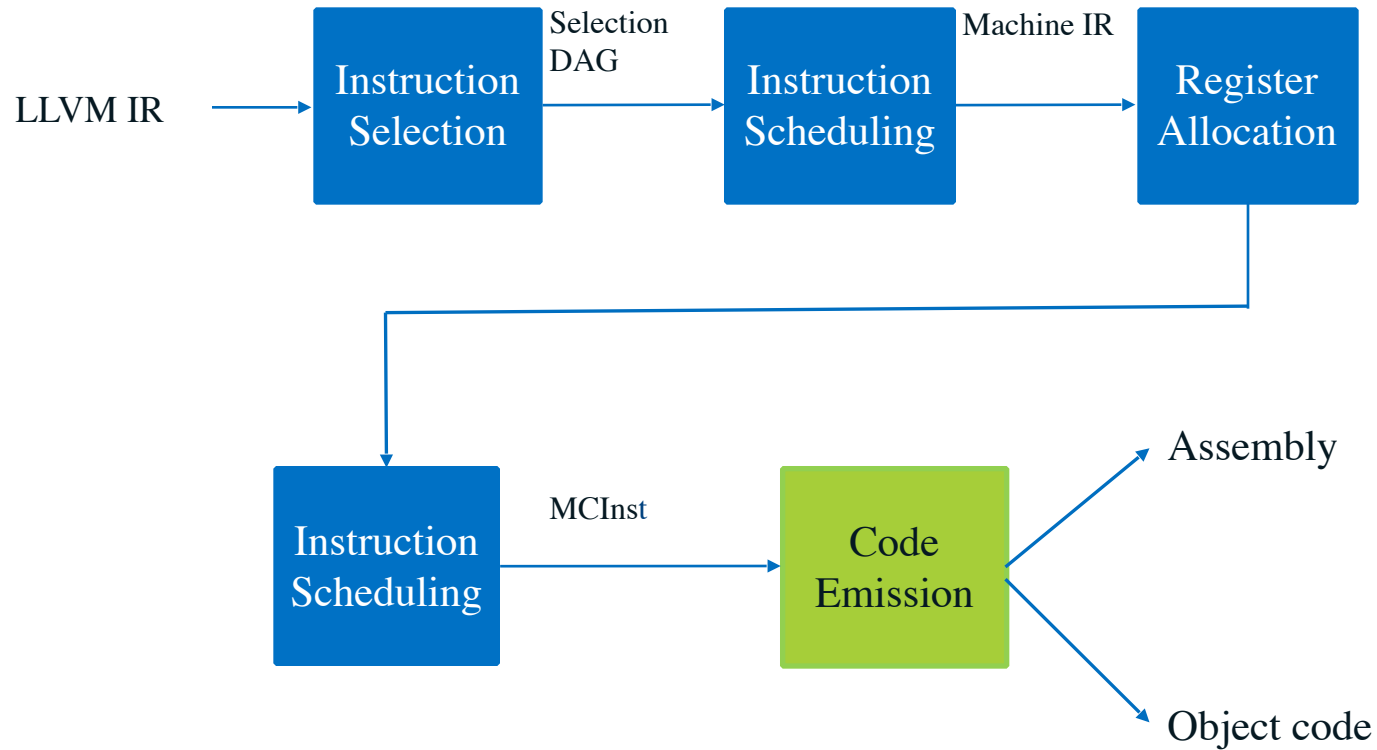
LLVM Backend Pipeline



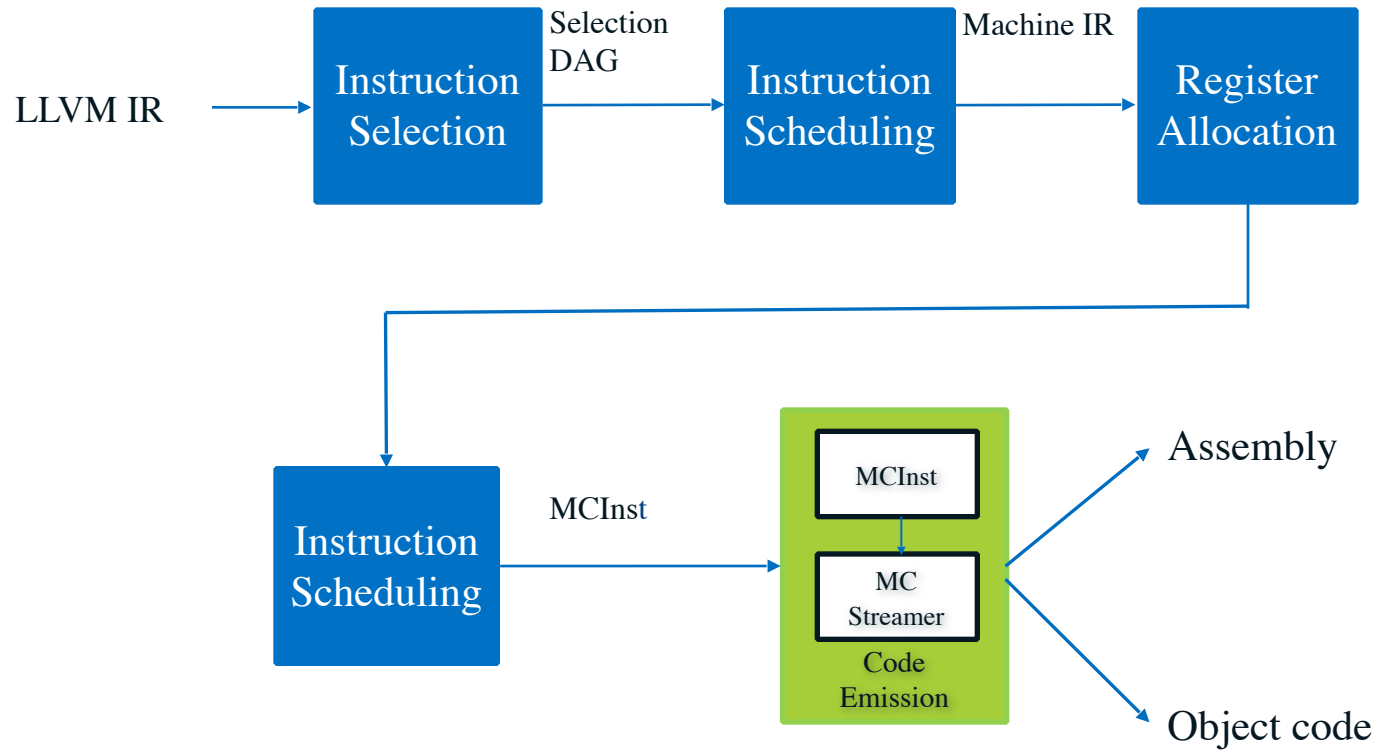
LLVM Backend Pipeline



LLVM Backend Pipeline



LLVM Backend Pipeline



Hands-on

Hands-On : Exercise 0

Build & install clang and llvm

- Objective
 - Learn how to download clang & llvm source from git
<https://github.com/jayakrish/llvm>
<https://github.com/jayakrish/clang>
or
<https://github.com/dekasthiti/llvm>
<https://github.com/dekasthiti/clang>
 - Learn how to configure and build clang & llvm
 - <https://github.com/jayakrish/Workshop-Materials>
 - <https://github.com/dekasthiti/Workshop-Materials>

Hands-On : Exercise 1

Hello World with clang

- Objective
 - Difference between clang, gcc, if any
 - Show clang error diagnostics
 - To show how to compile with clang.
 - How to use the different LLVM tools

CLANG vs GCC

CLANG	GCC
<i>ASTs and design are intended to be easily understandable</i>	<i>GCC has a very old codebase which presents a steep learning curve to new developers.</i>
<i>Clang is designed as an API from its inception, allowing it to be reused by source analysis tools, refactoring, IDEs (etc) as well as for code generation.</i>	<i>GCC is built as a monolithic static compiler, which makes it extremely difficult to use as an API and integrate into other tools.</i>
<i>Clang has none of these problems.</i>	<i>Various GCC design decisions make it very difficult to reuse: its build system is difficult to modify, you can't link multiple targets into one binary, you can't link multiple front-ends into one binary, it uses a custom garbage collector, uses global variables extensively, is not reentrant or multi-threadable, etc.</i>
<i>Clang does not implicitly simplify code as it parses it like GCC does. Doing so causes many problems for source analysis tools</i>	<i>If you write "x-x" in your source code, the GCC AST will contain "0", with no mention of 'x'. This is extremely bad for a refactoring tool that wants to rename 'x'</i>

CLANG vs GCC

CLANG	GCC
<i>Clang can serialize its AST out to disk and read it back into another program, which is useful for whole program analysis.</i>	<i>GCC does not have this.</i>
<i>Clang is much faster in compile time and uses far less memory than GCC.</i>	
<i>Clang aims to provide extremely clear and concise diagnostics (error and warning messages), and includes support for expressive diagnostics.</i>	<i>GCC's warnings are sometimes acceptable, but are often confusing and it does not support expressive diagnostics.</i>
<i>Clang uses a BSD license, which allows it to be embedded in software that is not GPL-licensed.</i>	<i>GCC is licensed under the GPL license</i>

LLVM IR Explained

MODULE – top level structure

global variables

functions

Basic Blocks – containers for instructions

library references

symbol-table

target characteristics

target datalayout: <datatype><size>:<abi>:<pref>

target triple:


<Architecture> - <sub> - <os> - <vendoe> - <abi>

Hands-On: Exercise 2

Generate LLVM IR using LLVM API

- Objective
 - Learn how to generate LLVM IR with LLVM API
 - Introduce LLVM concepts of Module, IRBuilder, Verifier, ExecutionEngine, PassManager
 - Explain how the LLVM API generator can be used to build a custom compiler.

Build Your Custom Compiler



Hands-On: Exercise 3

Print and view state and flow in a program

- Objective
 - Learn how to use LLVM tool opt, to view control flow graph
 - Learn how to use LLVM tool opt, to view call graph

Hands-On: Exercise 3

Debug — bash — 61x32

sdeka-mac01:Debug sdeka\$

I

```
libraryInfoWrapperPass>();  
LI() : nullptr;
```

In the worklist...

www LLVM LLVM Google Cour www LLVM x

HFA_enUS589US589&oq=llvm+instcombine&aqs=chrome..69i57....

Acquire and Release Memory Ordering for SF72

Search Sthiti

Search tools

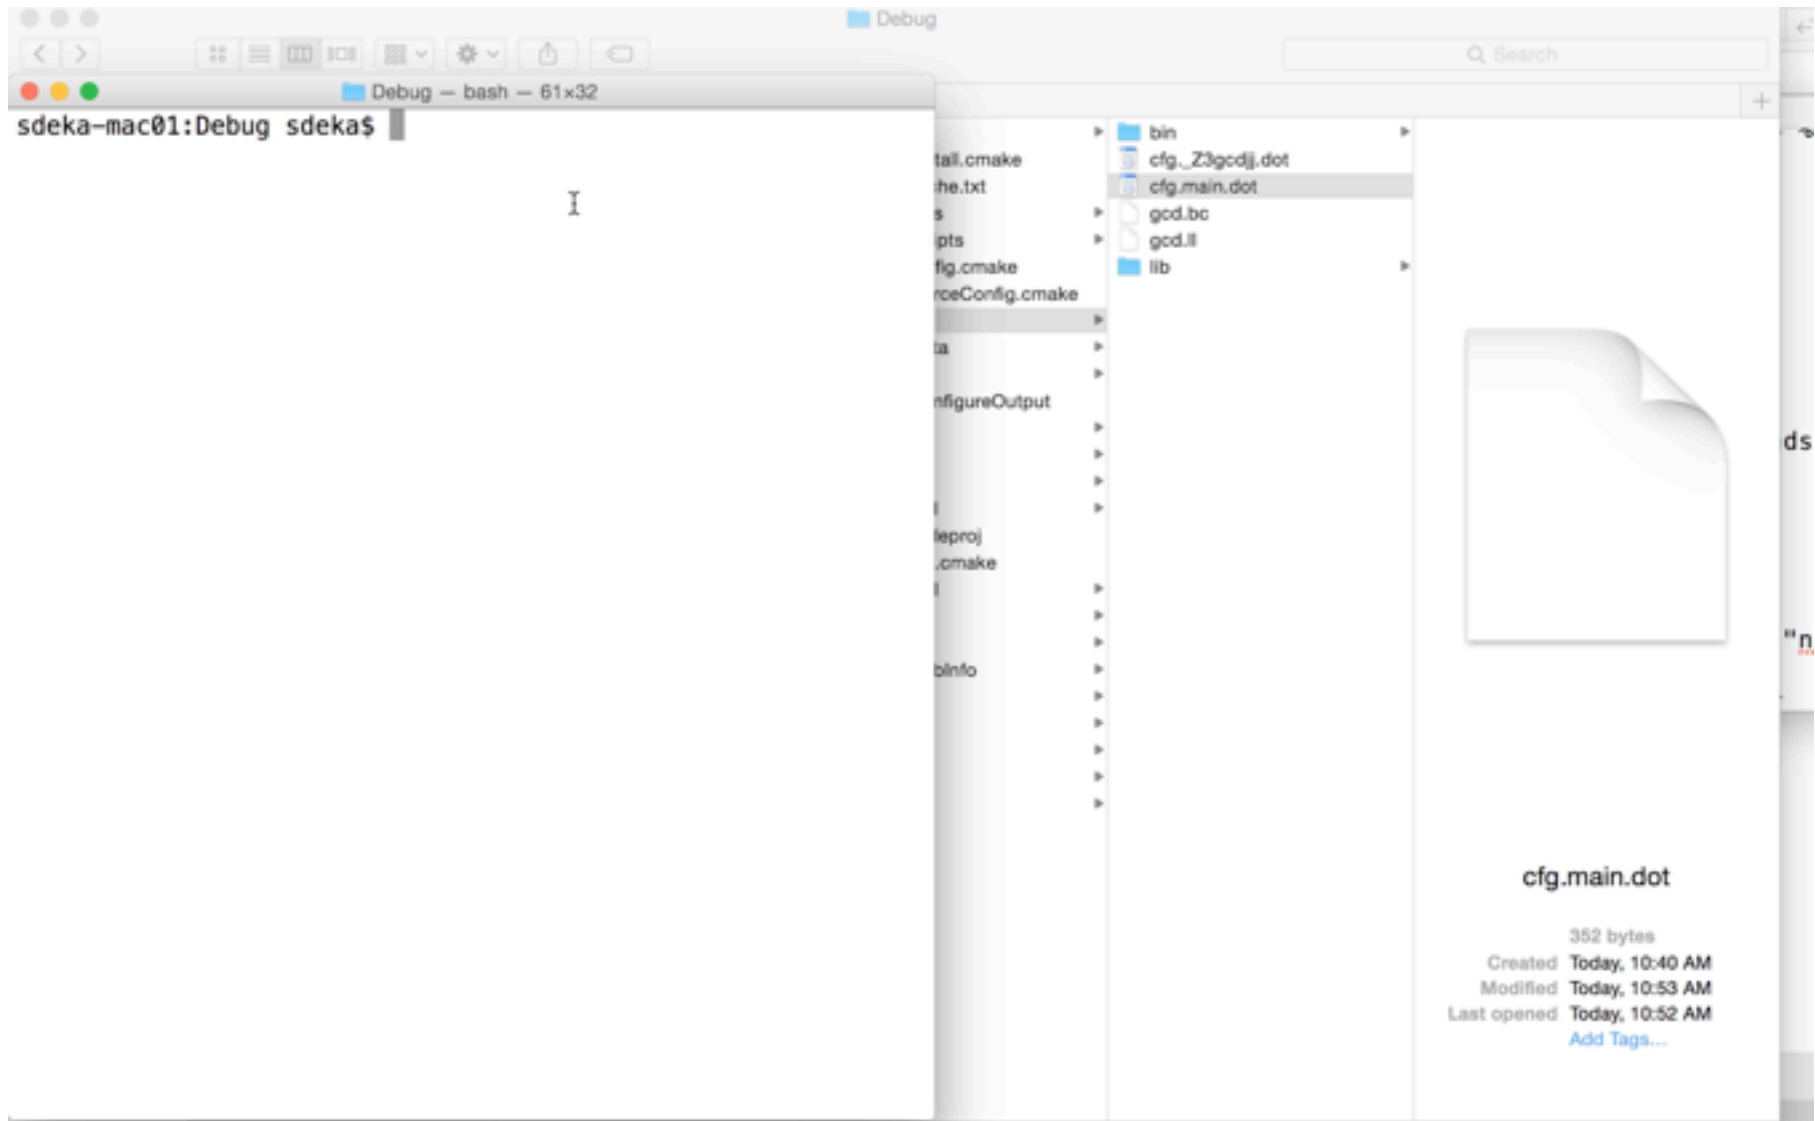
3.7 ...

bine checks for the
into ...

LLVM
00011 // This file
00012 // is ...

LLVM
bine ..
pp ...

Hands-On: Exercise 3



Hands-On: Exercise 4

Write a pass using LLVM API

- Objective
 - Learn how to write a pass
 - Learn about LLVM pass types
 - Learn about LLVM data structures such as iterators, dense maps
 - Learn about the tools, commands and headers used to do this exercise

Hands-On: Exercise 4

Pass in LLVM

- Pass Hierarchy

Hands-On: Exercise 4

Pass in LLVM

- Pass Hierarchy
 - Module Pass

Hands-On: Exercise 4

Pass in LLVM

- Pass Hierarchy
 - Module Pass
 - Function Pass

Hands-On: Exercise 4


Pass in LLVM

- Pass Hierarchy
 - Module Pass
 - Function Pass
 - Basic Block Pass

Hands-On: Exercise 4

Pass in LLVM


- Pass Hierarchy
 - Module Pass
 - Function Pass
 - Basic Block Pass



Hands-On: Exercise 4

Pass in LLVM


- Pass Examples



Hands-On: Exercise 4

Pass in LLVM


- Pass Examples
 - Module Pass
 - Eg: Call Graph



Hands-On: Exercise 4

Pass in LLVM


- Pass Examples
 - Module Pass
 - Eg: Call Graph
 - Function Pass
 - Eg: Constant Propagation



Hands-On: Exercise 4

Pass in LLVM


- Pass Examples
 - Module Pass
 - Eg: Call Graph
 - Function Pass
 - Eg: Constant Propagation
 - Basic Block Pass
 - Eg: Dead Instruction Elimination



Hands-On: Exercise 4

Pass in LLVM


- Pass Interfaces



Hands-On: Exercise 4

Pass in LLVM


- Pass Interfaces
 - Module Pass::
 - `runOnModule(Module &M)`



Hands-On: Exercise 4

Pass in LLVM


- Pass Interfaces
 - Module Pass::
 - `runOnModule(Module &M)`
 - Function Pass ::
 - `runOnFunction(Function &F)`



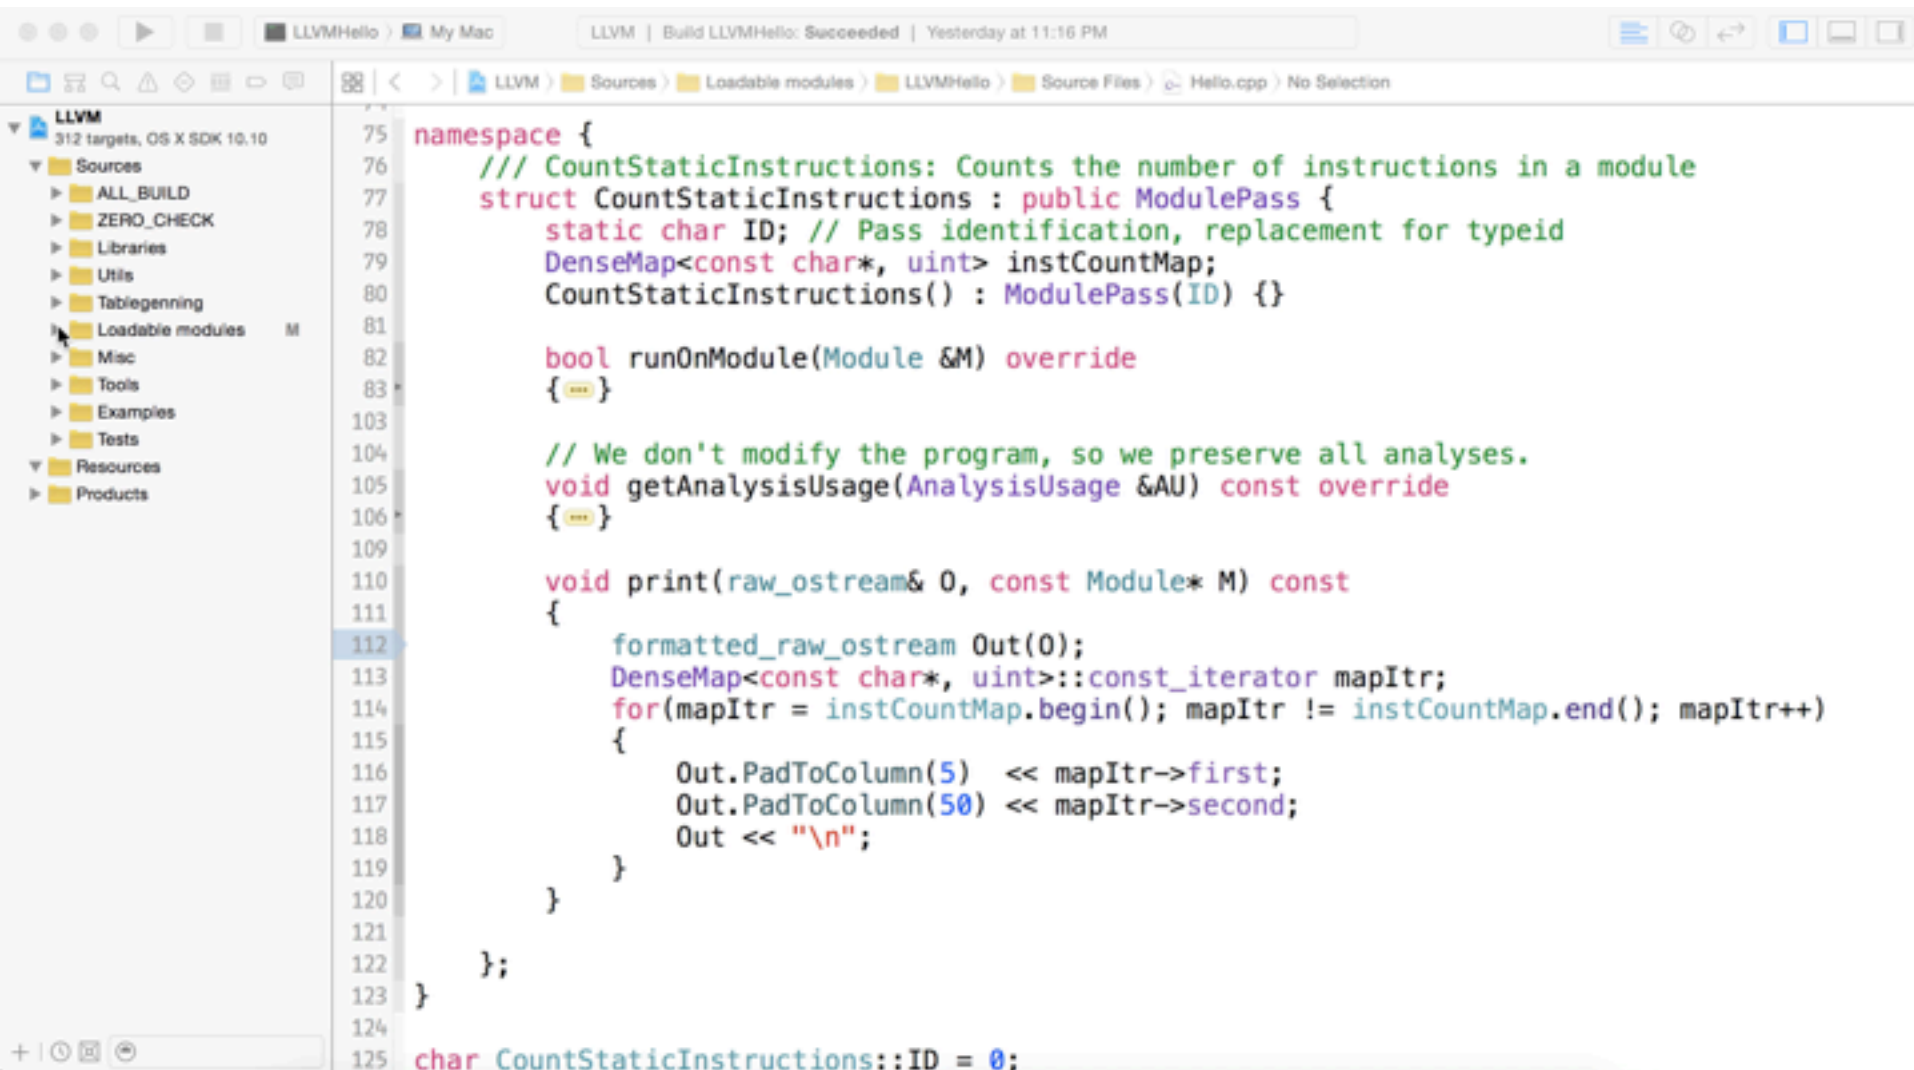
Hands-On: Exercise 4

Pass in LLVM

- Pass Interfaces
 - Module Pass::
 - `runOnModule(Module &M)`
 - Function Pass ::
 - `runOnFunction(Function &F)`
 - Basic Block Pass ::
 - `runOnBasicBlock(BasicBlock &B)`




Hands-On: Exercise 4 - Hello Pass

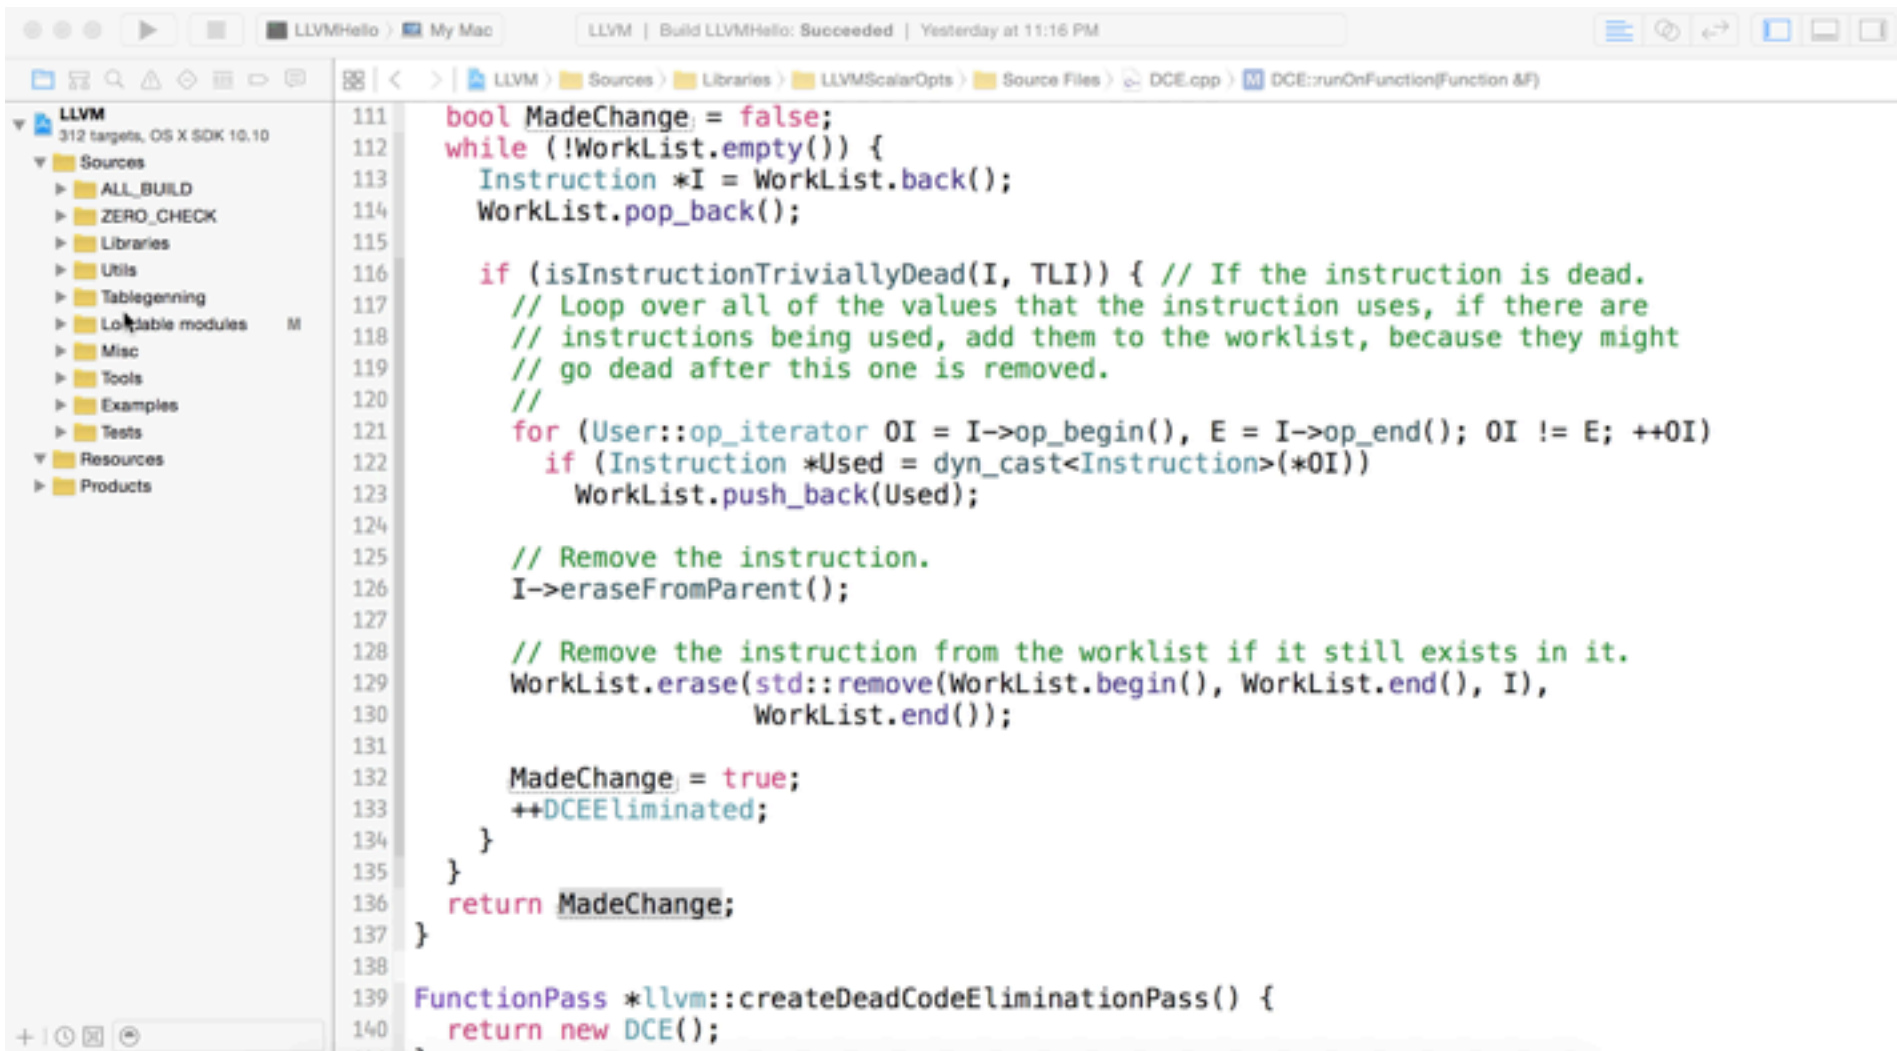


```
75 namespace {
76     /// CountStaticInstructions: Counts the number of instructions in a module
77     struct CountStaticInstructions : public ModulePass {
78         static char ID; // Pass identification, replacement for typeid
79         DenseMap<const char*, uint> instCountMap;
80         CountStaticInstructions() : ModulePass(ID) {}
81
82         bool runOnModule(Module &M) override
83         { ... }
84
85         // We don't modify the program, so we preserve all analyses.
86         void getAnalysisUsage(AnalysisUsage &AU) const override
87         { ... }
88
89         void print(raw_ostream& O, const Module* M) const
90         {
91             formatted_raw_ostream Out(O);
92             DenseMap<const char*, uint>::const_iterator mapItr;
93             for(mapItr = instCountMap.begin(); mapItr != instCountMap.end(); mapItr++)
94             {
95                 Out.PadToColumn(5) << mapItr->first;
96                 Out.PadToColumn(50) << mapItr->second;
97                 Out << "\n";
98             }
99         }
100     };
101 }
102
103 char CountStaticInstructions::ID = 0;
```

Hands-On: Exercise 4 - Hello Pass




Hands-On: Exercise 4 - Count Instructions



```
111 bool MadeChange = false;
112 while (!WorkList.empty()) {
113     Instruction *I = WorkList.back();
114     WorkList.pop_back();
115
116     if (isInstructionTriviallyDead(I, TLI)) { // If the instruction is dead.
117         // Loop over all of the values that the instruction uses, if there are
118         // instructions being used, add them to the worklist, because they might
119         // go dead after this one is removed.
120         //
121         for (User::op_iterator OI = I->op_begin(), E = I->op_end(); OI != E; ++OI)
122             if (Instruction *Used = dyn_cast<Instruction>(*OI))
123                 WorkList.push_back(Used);
124
125         // Remove the instruction.
126         I->eraseFromParent();
127
128         // Remove the instruction from the worklist if it still exists in it.
129         WorkList.erase(std::remove(WorkList.begin(), WorkList.end(), I),
130                        WorkList.end());
131
132         MadeChange = true;
133         ++DCEEliminated;
134     }
135 }
136 return MadeChange;
137 }
138
139 FunctionPass *llvm::createDeadCodeEliminationPass() {
140     return new DCE();
141 }
```

Hands-On: Exercise 4 - Count Instructions

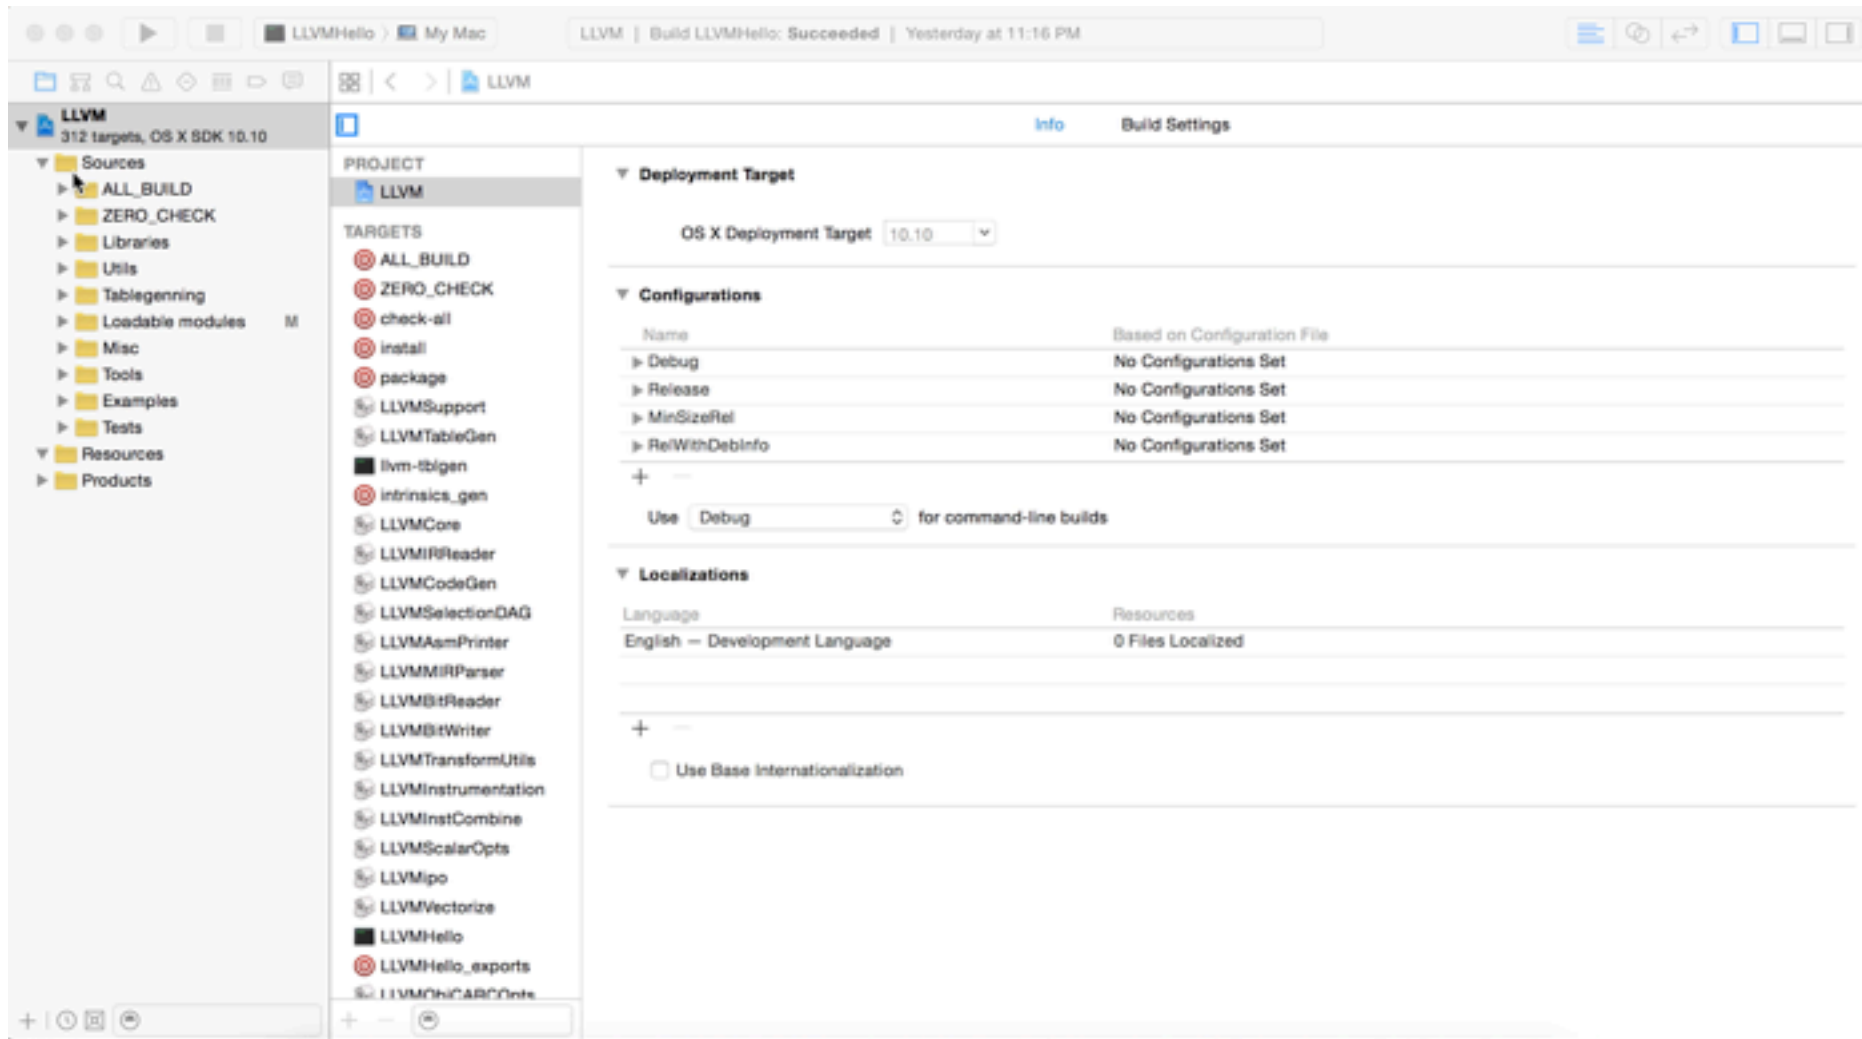


Hands-On: Exercise 5


Walk through an LLVM transformation pass

- Objective
 - Learn how to write an LLVM transformation pass
 - Learn how an LLVM transformation pass modifies the LLVM IR
 - Learn how to verify correctness of a transformation pass

Hands-On: Exercise 5 - Dead Code Elimination



Hands-On: Exercise 5 - Dead Code Elimination



Summing Up

- LLVM Compiler
 - Design overview
 - IR, API
 - Optimizations, Analysis Passes
- Tools
- Examples

References

<http://llvm.org/docs/index.html#>

<http://llvm.org/docs/GettingStarted.html>

http://clang.llvm.org/get_started.html

<http://llvm.org/docs/LangRef.html>


<http://cgo.org/cgo2015/conference/workshops-and-tutorials/#llvm>

<https://soco.intel.com/groups/inside-llvm>

BYOC!

Backup

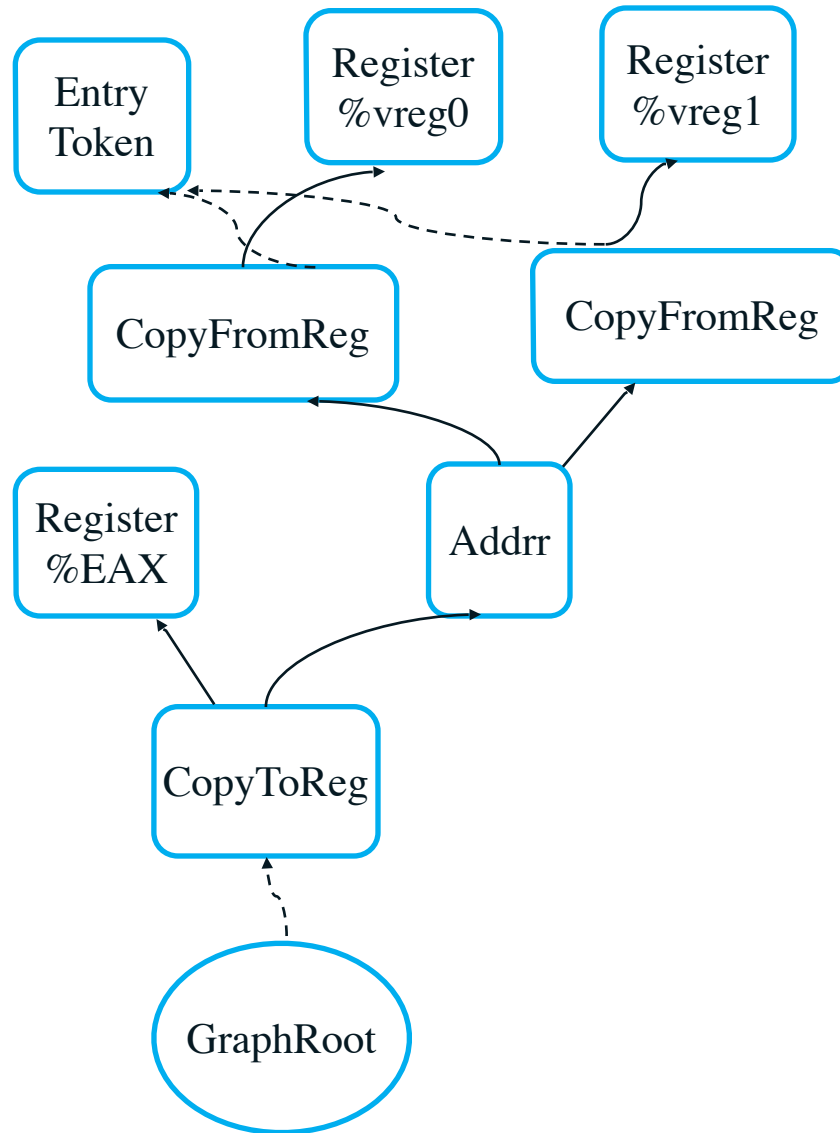
LLVM Backend - Instruction Scheduling



`$llc -print-machineinstrsmachine` dumps instructions after all registered passes

`$llc -print-machineinstrs=<pass-name>` dumps instructions after a specific pass

SelectionDAG Example



Example Machine IR

BB#0: derived from LLVM BB %entry

Live Ins: %I0 %I1

%vreg1<def> = COPY %I1; IntRegs:%vreg1

%vreg0<def> = COPY %I0; IntRegs:%vreg0

%vreg2<def> = ADDrr %vreg1, %vreg0; IntRegs:
%vreg2,%vreg1,%vreg0

%I0<def> = COPY %vreg2; IntRegs:%vreg2

RETL 8, %I0<imp-use>

SelectionDAG Instruction Select Phase Example

```
def MADDrr: Instx86<(outs GRRegs:$dst),  
    (ins GRRegs: $src1, GRRegs:$src2, GRRegs: $src3),  
    "madd $dst, $src1, $src2, $src3",  
    [(set i32:$dst, (add(mul(i32:$src1, i32:$src2), i32:$src3)))];
```

```
def ADDrr: Instx86<(outs GRRegs:$dst),  
    (ins GRRegs: $src1, GRRegs:$src2),  
    "add $dst, $src1, $src2",  
    [(set i32:$dst, (add i32:$src1, i32:$src2))];
```


MCInst Example

`addl -4(%rbp), %esi`

`## <MCInst #89 ADD32rm`

`## <MCOperand Reg:29>`

`## <MCOperand Reg:29>`

`## <MCOperand Reg:36>`

`## <MCOperand Imm:1>`

`## <MCOperand Reg:0>`

`## <MCOperand Imm:-4>`

`## <MCOperand Reg:0>>`

Assembly Example

```
.cfi_def_cfa_register %rbp
```

```
    movl    %edi, -4(%rbp)
```

```
    movl    %esi, -8(%rbp)
```

```
    addl    -4(%rbp), %esi
```

```
    movl    %esi, %eax
```

```
    popq    %rbp
```

```
    ret
```

```
.cfi_endproc
```

TableGen for LLVM backends

- LLVM's domain specific language
- Helps LLVM understand the target architecture
- Minimizes repetition and errors
- Declare machine specific aspects in single location but has multiple uses

Example:<Target>InstrInfo.td interpreted by AsmWriter backend or SelectionDAGISel backend

```
def ADDrr: Instx86<(outs GRRegs:$dst),  
  
    (ins GRRegs: $src1, GRRegs:$src2),  
  
    "add $dst, $src1, $src2",  
  
    [(set i32:$dst, (add i32:$src1, i32:$src2))];
```

Backend Code Structure

Main libraries are found in the lib directory and its subfolders

- CodeGen: generic code gen algorithms
- MC: Low level functionality for assemblers, disassemblers and object file types (eg: ELF COFF)
- TableGen: tool to generate C++ code based on high level descriptors found in .td file
- Target: Target specific implementation each within a different subfolder under Target eg: Target/x86

Generating DAG Graphs

```
llc -march=x86 -[OPTION] sum.ll
```

Where Option is:

- view-dag-combine1-dags displays the DAG after being built, before the first optimization pass.
- view-legalize-dags displays the DAG before Legalization.
- view-dag-combine2-dags displays the DAG before the second optimization pass.
- view-isel-dags displays the DAG before the Select phase.
- view-sched-dags displays the DAG before Scheduling.

LLVM Backend Classes

`SelectionDAGISel` – main base class for instruction selection

`SelectionDAGBuilder::visit` - `SelectionDAGISel` goes through IR and calls the `visit()` dispatcher on them

`TargetLowering` – an important interface to convey target specific information to target-independent algorithms

`SelectionDAGISel::Select` – custom instruction selection code is inserted here

`InstrEmitter::EmitMachineNode` – Scheduler uses this class to emit instructions into a `machinebasicblock`

TableGen Files

Target/<target>/<target>InstInfo.td – to implement instructions

Target/<target>/<target>RegisterInfo.td – to implement registers/ register classes

Target/<target>/<target>CallingConv.td – to implement calling conventions

<target>ISelLowering.cpp – Target specific hooks to lower instructions to target specific code. eg: LowerFormalArgs(), LowerRet()

<target>InstructionPrinter.cpp – printOperand() is given to the stream to print the instruction